

O'REILLY®

# Learning Ray

Flexible Distributed Python for Machine Learning



**Early  
Release**

Raw & Unedited

Compliments of



**anyscale**

Max Pumperla,  
Edward Oakes  
& Richard Liaw



---

# Learning Ray

*Flexible Distributed Python for Data Science*

With Early Release ebooks, you get books in their earliest form—the author’s raw and unedited content as they write—so you can take advantage of these technologies long before the official release of these titles.

*Max Pumperla, Edward Oakes, and Richard Liaw*

## **Learning Ray**

by Max Pumperla, Edward Oakes, and Richard Liaw

Copyright © 2023 Max Pumperla and O'Reilly Media inc. All rights reserved.

Printed in the United States of America.

Published by O'Reilly Media, Inc., 1005 Gravenstein Highway North, Sebastopol, CA 95472.

O'Reilly books may be purchased for educational, business, or sales promotional use. Online editions are also available for most titles ( <http://oreilly.com> ). For more information, contact our corporate/institutional sales department: 800-998-9938 or [corporate@oreilly.com](mailto:corporate@oreilly.com) .

**Editors:** Jeff Bleiel and Jessica Haberman

**Production Editor:** Katherine Tozer

**Interior Designer:** David Futato

**Cover Designer:** Karen Montgomery

**Illustrator:** Kate Dullea

April 2023: First Edition

### **Revision History for the Early Release**

2022-01-21: First Release

See <http://oreilly.com/catalog/errata.csp?isbn=9781098117221> for release details.

The O'Reilly logo is a registered trademark of O'Reilly Media, Inc. *Learning Ray*, the cover image, and related trade dress are trademarks of O'Reilly Media, Inc.

The views expressed in this work are those of the authors, and do not represent the publisher's views. While the publisher and the authors have used good faith efforts to ensure that the information and instructions contained in this work are accurate, the publisher and the authors disclaim all responsibility for errors or omissions, including without limitation responsibility for damages resulting from the use of or reliance on this work. Use of the information and instructions contained in this work is at your own risk. If any code samples or other technology this work contains or describes is subject to open source licenses or the intellectual property rights of others, it is your responsibility to ensure that your use thereof complies with such licenses and/or rights.

---

# Table of Contents

<b>1. An Overview of Ray.....</b>	<b>5</b>
What is Ray?	6
What led to Ray?	6
Three Layers: Core, Libraries & Ecosystem	9
A Distributed Computing Framework	9
A Suite of Data Science Libraries	12
Machine Learning and the Data Science Workflow	12
Data Processing with Ray Data	14
Model Training	16
Hyperparameter Tuning	20
Model Serving	22
A Growing Ecosystem	24
How Ray Integrates and Extends	24
Ray as Distributed Interface	24
Summary	25
<b>2. Getting Started With Ray Core.....</b>	<b>27</b>
An Introduction To Ray Core	27
A First Example Using the Ray API	28
An Overview of the Ray Core API	38
Design Principles	39
Understanding Ray System Components	40
Scheduling and Executing Work on a Node	40
The Head Node	43
Distributed Scheduling and Execution	44
Summary	46

<b>3. Building Your First Distributed Application With Ray Core.....</b>	<b>47</b>
Creating A Ray Core App	48
A Simple Maze Problem	48
Building a Simulation	53
Training a Reinforcement Learning Model	56
Building a Distributed Ray Trainer	60
Summary	64

---

# An Overview of Ray

## A Note for Early Release Readers

With Early Release ebooks, you get books in their earliest form—the author’s raw and unedited content as they write—so you can take advantage of these technologies long before the official release of these titles.

One of the reasons we need efficient distributed computing is that we’re collecting ever more data with a large variety at increasing speeds. The storage systems, data processing and analytics engines that have emerged in the last decade are crucially important to the success of many companies. Interestingly, most “big data” technologies are built for and operated by (data) engineers, that are in charge of data collection and processing tasks. The rationale is to free up data scientists to do what they’re best at. As a data science practitioner you might want to focus on training complex machine learning models, running efficient hyperparameter selection, building entirely new and custom models or simulations, or serving your models to showcase them. At the same time you simply might *have to* scale them to a compute cluster. To do that, the distributed system of your choice needs to support all of these fine-grained “big compute” tasks, potentially on specialized hardware. Ideally, it also fits into the big data tool chain you’re using and is fast enough to meet your latency requirements. In other words, distributed computing has to be powerful and flexible enough for complex data science workloads — and Ray can help you with that.

Python is likely the most popular language for data science today, and it’s certainly the one I find the most useful for my daily work. By now it’s over 30 years old, but has a still growing and active community. The rich [PyData ecosystem](#) is an essential part of a data scientist’s toolbox. How can you make sure to scale out your workloads while still leveraging the tools you need? That’s a difficult problem, especially since

communities can't be forced to just toss their toolbox, or programming language. That means distributed computing tools for data science have to be built for their existing community.

## What is Ray?

What I like about Ray is that it checks all the above boxes. It's a flexible distributed computing framework build for the Python data science community. Ray is easy to get started and keeps simple things simple. Its core API is as lean as it gets and helps you reason effectively about the distributed programs you want to write. You can efficiently parallelize Python programs on your laptop, and run the code you tested locally on a cluster practically without any changes. Its high-level libraries are easy to configure and can seamlessly be used together. Some of them, like Ray's reinforcement learning library, would have a bright future as standalone projects, distributed or not. While Ray's core is built in C++, it's been a Python-first framework since day one, integrates with many important data science tools, and can count on a growing ecosystem.

Distributed Python is not new, and Ray is not the first framework in this space (nor will it be the last), but it is special in what it has to offer. Ray is particularly strong when you combine several of its modules and have custom, machine learning heavy workloads that would be difficult to implement otherwise. It makes distributed computing easy enough to run your complex workloads flexibly by leveraging the Python tools you know and want to use. In other words, by *learning Ray* you get to know *flexible distributed Python for data science*.

In this chapter you'll get a first glimpse at what Ray can do for you. We will discuss the three layers that make up Ray, namely its core engine, its high-level libraries and its ecosystem. Throughout the chapter we'll show you first code examples to give you a feel for Ray, but we defer any in-depth treatment of Ray's APIs and components to later chapters. You can view this chapter as an overview of the whole book as well.

## What led to Ray?

Programming distributed systems is hard. It requires specific knowledge and experience you might not have. Ideally, such systems get out of your way and provide abstractions to let you focus on your job. But in practice “all non-trivial abstractions, to some degree, are leaky” ([Spolsky](#)), and getting clusters of computers to do what you want is undoubtedly difficult. Many software systems require resources that far exceed what single servers can do. Even if one server was enough, modern systems need to be failsafe and provide features like high availability. That means your applications might have to run on multiple machines, or even datacenters, just to make sure they're running reliably.



Even if you're not too familiar with machine learning (ML) or more generally artificial intelligence (AI) as such, you must have heard of recent breakthroughs in the field. To name just two, systems like [Deepmind's AlphaFold](#) for solving the protein folding problem, or [OpenAI's Codex](#) that's helping software developers with the tedious parts of their job, have made the news lately. You might also have heard that ML systems generally require large amounts of data to be trained. OpenAI has shown exponential growth in compute needed to train AI models in their paper "[AI and Compute](#)". The operations needed for AI systems in their study is measured in peta-flops (thousands of trillion operations per second), and has been *doubling every 3.4 months* since 2012.

Compare this to Moore's Law<sup>1</sup>, which states that the number of transistors in computers would double every two years. Even if you're bullish on Moore's law, you can see how there's a clear need for distributed computing in ML. You should also understand that many tasks in ML can be naturally decomposed to run in parallel. So, why not speed things up if you can?

Distributed computing is generally perceived as hard. But why is that? Shouldn't it be realistic to find good abstractions to run your code on clusters, without having to constantly think about individual machines and how they interoperate? What if we specifically focused on AI workloads?

Researchers at [RISELab](#) at UC Berkeley created Ray to address these questions. None of the tools existing at the time met their needs. They were looking for easy ways to speed up their workloads by distributing them to compute clusters. The workloads they had in mind were quite flexible in nature and didn't fit into the analytics engines available. At the same time, RISELab wanted to build a system that took care of how the work was distributed. With reasonable default behaviors in place, researchers should be able to focus on their work. And ideally they should have access to all their favorite tools in Python. For this reason, Ray was built with an emphasis on high-performance and heterogeneous workloads. [Anyscale](#), the company behind Ray, is building a managed Ray Platform and offers hosted solutions for your Ray applications. Let's have a look at an example of what kinds of applications Ray was designed for.

## Flexible Workloads in Python & Reinforcement Learning

One of my favorite apps on my phone can automatically classify or "label" individual plants in our garden. It works by simply showing it a picture of the plant in question. That's immensely helpful, as I'm terrible at distinguishing them all. (I'm not bragging

---

<sup>1</sup> Moore's Law held for a long time, but there might be signs that it's slowing down. We're not here to argue it, though. What's important is not that our computers generally keep getting faster, but the relation to the amount of compute we need.

about the size of my garden, I'm just bad at it.) In the last couple of years we've seen a surge of impressive applications like that.

Ultimately, the promise of AI is to build intelligent agents that go far beyond classifying objects. Imagine an AI application that not only knows your plants, but can take care of to them, too. Such an application would have to

- operate in dynamic environments (like the change of seasons),
- react to changes in the environment (like a heavy storm or pests attacking your plants),
- take sequences of actions (like watering and fertilizing plants),
- and accomplish long-term goals (like prioritizing plant health).

By observing its environment such an AI would also learn to explore the possible actions it could take and come up with better solutions over time. If you feel like this example is artificial or too far out, it's not difficult to come up with examples on your own that share all the above requirements. Think of managing and optimizing a supply chain, strategically restocking a warehouse considering fluctuating demands, or orchestrating the processing steps in an assembly line. Another famous example of what you could expect from an AI would be Stephen Wozniak's famous "Coffee Test". If you're invited to a friend's house, you can navigate to the kitchen, spot the coffee machine and all necessary ingredients, figure out how to brew a cup of coffee, and sit down to enjoy it. A machine should be able to do the same, except the last part might be a bit of a stretch. What other examples can you think of?

You can frame all the above requirements naturally in a subfield of machine learning called reinforcement learning (RL). We've dedicated all of ??? to RL. For now, it's enough to understand that it's about agents interacting with their environment by observing it and emitting actions. In RL, agents evaluate their environments by attributing a reward (e.g., how healthy is my plant on a scale from 1 to 10). The term "reinforcement" comes from the fact that agents will hopefully learn to seek out behaviour that leads to good outcomes (high reward), and shy away from punishing situations (low or negative reward). The interaction of agents with their environment is usually modeled by creating a computer simulation of it. These simulations can become complicated quite quickly, as you might imagine from the examples we've given.

We don't have gardening robots like the one I've sketched yet. And we don't know which AI paradigm will get us there.<sup>2</sup> What I do know is that the world is full of complex, dynamic and interesting examples that we need to tackle. For that we need computational frameworks that help us do that, and Ray was built to do exactly that.

---

<sup>2</sup> For the experts among you, I don't claim that RL is the answer. RL is just a paradigm that naturally fits into this discussion of AI goals.

RISELab created Ray to build and run complex AI applications at scale, and reinforcement learning has been an integral part of Ray from the start.

## Three Layers: Core, Libraries & Ecosystem

Now that you know why Ray was built and what its creators had in mind, let's look at the three layers of Ray.

1. A low-level, distributed computing framework for Python with a concise core API.<sup>3</sup>
2. A set of high-level libraries for data science built and maintained by the creators of Ray.
3. A growing ecosystem of integrations and partnerships with other notable projects.

There's a lot to unpack here, and we'll look into each of these layers individually in the remainder of this chapter. You can imagine Ray's core engine with its API at the center of things, on which everything else builds. Ray's data science libraries build on top of it. In practice, most data scientists will use these higher level libraries directly and won't often need to resort to the core API. The growing number of third-party integrations for Ray is another great entrypoint for experienced practitioners. Let's look into each one of the layers one by one.

## A Distributed Computing Framework

At its core, Ray is a distributed computing framework. We'll provide you with just the basic terminology here, and talk about Ray's architecture in depth in [Chapter 2](#). In short, Ray sets up and manages clusters of computers so that you can run distributed tasks on them. A ray cluster consists of nodes that are connected to each other via a network. You program against the so-called *driver*, the program root, which lives on the *head node*. The driver can run *jobs*, that is a collection of tasks, that are run on the nodes in the cluster. Specifically, the individual tasks of a job are run on *worker* processes on *worker nodes*. [Figure 1-1](#) illustrates the basic structure of a Ray cluster.

---

<sup>3</sup> This is a Python book, so we'll exclusively focus on it. But you should at least know that Ray also has a Java API, which at this point is less mature than its Python equivalent.




Figure 1-1. The basic components of a Ray cluster.

What's interesting is that a Ray cluster can also be a *local cluster*, i.e. a cluster consisting just of your own computer. In this case, there's just one node, namely the head node, which has the driver process and some worker processes. The default number of worker processes is the number of CPUs available on your machine.

With that knowledge at hand, it's time to get your hands dirty and run your first local Ray cluster. Installing Ray<sup>4</sup> on any of the major operating systems should work seamlessly using `pip`:

```
pip install "ray[rllib, serve, tune]"==1.9.0
```

With a simple `pip install ray` you would have installed just the very basics of Ray. Since we want to explore some advanced features, we installed the “extras” `rllib`, `serve` and `tune`, which we'll discuss in a bit. Depending on your system configuration you may not need the quotation marks in the above installation command.

Next, go ahead and start a Python session. You could use the `ipython` interpreter, which I find to be the most suitable environment for following along simple examples. If you don't feel like typing in the commands yourself, you can also jump into the [jupyter notebook for this chapter](#) and run the code there. The choice is up to you, but in any case please remember to use Python version 3.7 or later. In your Python session you can now easily import and initialize Ray as follows:

---

<sup>4</sup> We're using Ray version 1.9.0 at this point, as it's the latest version available as of this writing.

*Example 1-1.*

```
import ray
ray.init()
```

With those two lines of code you’ve started a Ray cluster on your local machine. This cluster can utilize all the cores available on your computer as workers. In this case you didn’t provide any arguments to the `init` function. If you wanted to run Ray on a “real” cluster, you’d have to pass more arguments to `init`. The rest of your code would stay the same.

After running this code you should see output of the following form (we use ellipses to remove the clutter):

```
... INFO services.py:1263 -- View the Ray dashboard at http://127.0.0.1:8265
{'node_ip_address': '192.168.1.41',
 'raylet_ip_address': '192.168.1.41',
 'redis_address': '192.168.1.41:6379',
 'object_store_address': '.../sockets/plasma_store',
 'raylet_socket_name': '.../sockets/raylet',
 'webui_url': '127.0.0.1:8265',
 'session_dir': '...',
 'metrics_export_port': 61794,
 'node_id': '...'}
```

This indicates that your Ray cluster is up and running. As you can see from the first line of the output, Ray comes with its own, pre-packaged dashboard. In all likelihood you can check it out at <http://127.0.0.1:8265>, unless your output shows a different port. If you want you can take your time to explore the dashboard for a little. For instance, you should see all your CPU cores listed and the total utilization of your (trivial) Ray application. We’ll come back to the dashboard in later chapters.

We’re not quite ready to dive into all the details of a Ray cluster here. To jump ahead just a little, you might see the `raylet_ip_address`, which is a reference to a so-called *Raylet*, which is responsible for scheduling tasks on your worker nodes. Each Raylet has a store for distributed objects, which is hinted at by the `object_store_address` above. Once tasks are scheduled, they get executed by worker processes. In [Chapter 2](#) you’ll get a much better understanding of all these components and how they make up a Ray cluster.

Before moving on, we should also briefly mention that the Ray core API is very accessible and easy to use. But since it is also a rather low-level interface, it takes time to build interesting examples with it. [Chapter 2](#) has an extensive first example to get you started with the Ray core API, and in [Chapter 3](#) you’ll see how to build a more interesting Ray application for reinforcement learning.

Right now your Ray cluster doesn't do much, but that's about to change. After giving you a quick introduction to the data science workflow in the following section, you'll run your first concrete Ray examples.

## A Suite of Data Science Libraries

Moving on to the second layer of Ray, in this section we'll introduce all the data science libraries that Ray comes with. To do so, let's first take a bird's eye view on what it means to do data science. Once you understand this context, it's much easier to place Ray's higher-level libraries and see how they can be useful to you. If you have a good idea of the data science process, you can safely skip ahead to section [“Data Processing with Ray Data” on page 14](#).

## Machine Learning and the Data Science Workflow

The somewhat elusive term “data science” (DS) evolved quite a bit in recent years, and you can find many definitions of varying usefulness online.<sup>5</sup> To me, it's *the practice of gaining insights and building real-world applications by leveraging data*. That's quite a broad definition, and you don't have to agree with me. My point is that data science is an inherently practical and applied field that centers around building and understanding things, which makes fairly little sense in a *purely* academic context. In that sense, describing practitioners of this field as “data scientists” is about as bad of a misnomer as describing hackers as “computer scientists”.<sup>6</sup>

Since you are familiar with Python and hopefully bring a certain craftsmanship attitude with you, we can approach the Ray's data science libraries from a very pragmatic angle. Doing data science in practice is an iterative process that goes something like this:

- Requirements engineering: You talk to stakeholders to identify the problems you need to solve and clarify the requirements for this project.
- Data collection: Then you source, collect and inspect the data.
- Data processing: Afterwards you process the data such that you can tackle the problem.

---

<sup>5</sup> I never liked the categorization of data science as an intersection of disciplines, like maths, coding and business. Ultimately, that doesn't tell you what practitioners *do*. It doesn't do a cook justice to tell them they sit at the intersection of agriculture, thermodynamics and human relations. It's not wrong, but also not very helpful.

<sup>6</sup> As a fun exercise, I recommend reading Paul Graham's famous [“Hackers and Painters”](#) essay on this topic and replace “computer science” with “data science”. What would hacking 2.0 be?

- **Model building:** You then move on to build a model (in the broadest sense) using the data. That could be a dashboard with important metrics, a visualisation, or a machine learning model, among many other things.
- **Model evaluation:** The next step is to evaluate your model against the requirements in the first step.
- **Deployment:** If all goes well (it likely doesn't), you deploy your solution in a production environment. You should understand this as an ongoing process that needs to be monitored, not as a one-off step.
- Otherwise you need to circle back and start from the top. The most likely outcome is that you need to improve your solution in various ways, even after initial deployment.

Machine learning is not necessarily part of this process, but you can see how building smart applications or gaining insights might benefit from ML. Building a face detection app into your social media platform, for better or worse, might be one example of that. When the data science process just described explicitly involves building machine learning models, you can further specify some steps:

- *Data processing:* To train machine learning models, you need data in a format that is understood by your ML model. The process of transforming and selecting what data should be fed into your model is often called *feature engineering*. This step can be messy. You'll benefit a lot if you can rely on common tools to do the job.
- *Model training:* In ML you need to train your algorithms on data that got processed in the last step. This includes selecting the right algorithm for the job, and it helps if you can choose from a wide variety.
- *Hyperparameter tuning:* Machine learning models have parameters that are tuned in the model training step. Most ML models also have another set of parameters, called *hyperparameters* that can be modified prior to training. These parameters can heavily influence the performance of your resulting ML model and need to be tuned properly. There are good tools to help automate that process.
- *Model serving:* Trained models need to be deployed. To serve a model means to make it available to whoever needs access by whatever means necessary. In prototypes, you often use simple HTTP servers, but there are many specialised software packages for ML model serving.

This list is by no means exhaustive. Don't worry if you've never gone through these steps or struggle with the terminology, we'll come back to this in much more detail in later chapters. If you want to understand more about the holistic view of the data science process when building machine learning applications, the book [Building Machine Learning Powered Applications](#) is dedicated to it entirely.

Figure 1-2 gives an overview of the steps we just discussed:

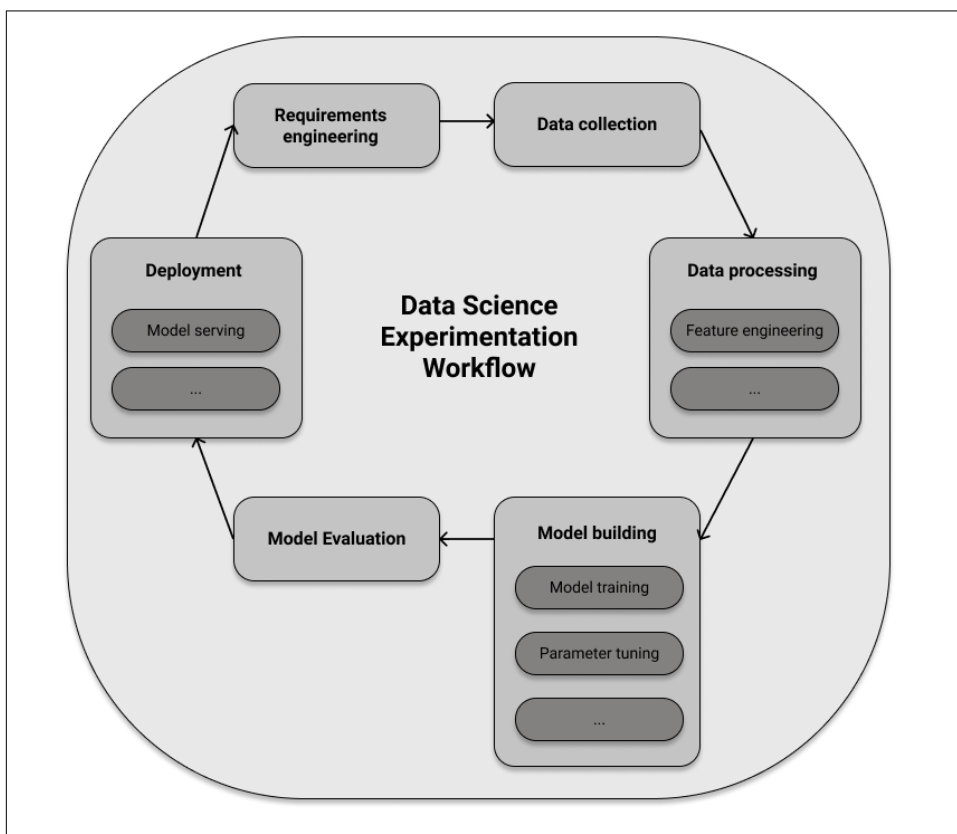


Figure 1-2. An overview of the data science experimentation workflow using machine learning.

At this point you might be wondering how any of this relates to Ray. The good news is that Ray has a dedicated library for each of the four ML-specific tasks above, covering data processing, model training, hyperparameter tuning and model serving. And the way Ray is designed, all these libraries are *distributed by construction*. Let's walk through each of them one-by-one.

## Data Processing with Ray Data

The first high-level library of Ray we talk about is called “Ray Data”. This library contains a data structure aptly called `Dataset`, a multitude of connectors for loading data from various formats and systems, an API for transforming such datasets, a way to build data processing pipelines with them, and many integrations with other data



processing frameworks. The Dataset abstraction builds on the powerful [Arrow framework](#).

To use Ray Data, you need to install Arrow for Python, for instance by running `pip install pyarrow`. We'll now discuss a simple example that creates a distributed Dataset on your local Ray cluster from a Python data structure. Specifically, you'll create a dataset from a Python dictionary containing a string name and an integer-valued data for 10000 entries:

*Example 1-2.*

```
import ray

items = [{"name": str(i), "data": i} for i in range(10000)]
ds = ray.data.from_items(items) ❶
ds.show(5) ❷
```

❶ Creating a Dataset by using `from_items` from the `ray.data` module.

❷ Printing the first 10 items of the Dataset.

To show a Dataset means to print some of its values. You should see precisely 5 so-called `ArrowRow` elements on your command line, like this:

```
ArrowRow({'name': '0', 'data': 0})
ArrowRow({'name': '1', 'data': 1})
ArrowRow({'name': '2', 'data': 2})
ArrowRow({'name': '3', 'data': 3})
ArrowRow({'name': '4', 'data': 4})
```

Great, now you have some distributed rows, but what can you do with that data? The Dataset API bets heavily on functional programming, as it is very well suited for data transformations. Even though Python 3 made a point of hiding some of its functional programming capabilities, you're probably familiar with functionality such as `map`, `filter` and others. If not, it's easy enough to pick up. `map` takes each element of your dataset and transforms it into something else, in parallel. `filter` removes data points according to a boolean filter function. And the slightly more elaborate `flat_map` first maps values similarly to `map`, but then also “flattens” the result. For instance, if `map` would produce a list of lists, `flat_map` would flatten out the nested lists and give you just a list. Equipped with these three functional API calls, let's see how easily you can transform your dataset `ds`:

*Example 1-3. Transforming a Dataset with common functional programming routines.*

```
squares = ds.map(lambda x: x["data"] ** 2) ❶
```

```
evens = squares.filter(lambda x: x % 2 == 0) ❷
evens.count()

cubes = evens.flat_map(lambda x: [x, x**3]) ❸
sample = cubes.take(10) ❹
print(sample)
```

- ❶ We map each row of `ds` to only keep the square value of its data entry.
- ❷ Then we filter the squares to only keep even numbers (a total of 5000 elements).
- ❸ We then use `flat_map` to augment the remaining values with their respective cubes.
- ❹ To take a total of 10 values means to leave Ray and return a Python list with these values that we can print.

The drawback of Dataset transformations is that each step gets executed synchronously. In example [Example 1-3](#) this is a non-issue, but for complex tasks that e.g. mix reading files and processing data, you want an execution that can overlap individual tasks. `DatasetPipeline` does exactly that. Let's rewrite the last example into a pipeline.

*Example 1-4.*

```
pipe = ds.window() ❶
result = pipe\
    .map(lambda x: x["data"] ** 2)\
    .filter(lambda x: x % 2 == 0)\
    .flat_map(lambda x: [x, x**3]) ❷
result.show(10)
```

- ❶ You can turn a Dataset into a pipeline by calling `.window()` on it.
- ❷ Pipeline steps can be chained to yield the same result as before.

There's a lot more to be said about Ray Data, especially its integration with notable data processing systems, but we'll have to defer an in-depth discussion until ???.

## Model Training


Moving on to the next set of libraries, let's look at the distributed training capabilities of Ray. For that, you have access to two libraries. One is dedicated to reinforcement learning specifically, the other one has a different scope and is aimed primarily at supervised learning tasks.

## Reinforcement Learning with Ray RLlib

Let's start with *Ray RLlib* for reinforcement learning. This library is powered by the modern ML frameworks TensorFlow and PyTorch, and you can choose which one to use. Both frameworks seem to converge more and more conceptually, so you can pick the one you like most without losing much in the process. Throughout the book we use TensorFlow for consistency. Go ahead and install it with `pip install tensorflow` right now.

One of the easiest ways to run examples with RLlib is to use the command line tool `rllib`, which we've already implicitly installed earlier with `pip`. Once you run more complex examples in `???`, you will mostly rely on its Python API, but for now we just want to get a first taste of running RL experiments.

We'll look at a fairly classical control problem of balancing a pendulum. Imagine you have a pendulum like the one in figure [Figure 1-3](#), fixed at a single point and subject to gravity. You can manipulate that pendulum by giving it a push from the left or the right. If you assert just the right amount of force, the pendulum might remain in an upright position. That's our goal - and the question is whether we can teach a reinforcement learning algorithm to do so for us.



*Figure 1-3. Controlling a simple pendulum by asserting force to the left or the right.*

Specifically, we want to train a reinforcement learning agent that can push to the left or right, thereby acting on its environment (manipulating the pendulum) to reach the

“upright position” goal for which it will be rewarded. To tackle this problem with Ray RLlib, store the following content in a file called `pendulum.yml`.

*Example 1-5.*

```
pendulumppo:
  env: Pendulumv0 ❶
  run: PPO ❷
  checkpoint_freq: 5 ❸
  stop:
    episode_reward_mean: 800 ❹
  config:
    lambda: 0.1 ❺
    gamma: 0.95
    lr: 0.0003
    num_sgd_iter: 6
```

- ❶ The Pendulum-v0 environment simulates the pendulum problem we just described.
- ❷ We use a powerful RL algorithm called Proximal Policy Optimization, or PPO.
- ❸ After every five “training iterations” we checkpoint a model.
- ❹ Once we reach a reward of -800, we stop the experiment.
- ❺ The PPO needs some RL-specific configuration to make it work for this problem.

The details of this configuration file don’t matter much at this point, don’t get distracted by them. The important part is that you specify the built-in Pendulum-v0 environment and sufficient RL-specific configuration to ensure the training procedure works. The configuration is a simplified version of one of Ray’s **tuned examples**. We chose this one because it doesn’t require any special hardware and finishes in a matter of minutes. If your computer is powerful enough, you can try to run the tuned example as well, which should yield much better results. To train this pendulum example you can now simply run:

```
rllib train -f pendulum.yml
```

If you want, you can check the output of this Ray program and see how the different metrics evolve during the training procedure. In case you don’t want to create this file on your own, and want to run an experiment which gives you much better results, you can also run this:

```
curl https://raw.githubusercontent.com/maxpumperla/learning_ray/main/notebooks/pendulum.yml -o per
rllib train -f pendulum.yml
```

In any case, assuming the training program finished, we can now check how well it worked. To visualize the trained pendulum you need to install one more Python library with `pip install pygame`. The only other thing you need to figure out is where Ray stored your training progress. When you run `rllib train` for an experiment, Ray will create a unique experiment ID for you and stores results in a subfolder of `~/ray-results` by default. For the training configuration we used, you should see a folder with results that looks like `~/ray_results/pendulum-ppo/PP0_Pendulum-v0_<experiment_id>`. During the training procedure intermediate model checkpoints get generated in the same folder. For instance, I have a folder on my machine called:

```
~/ray_results/pendulum-ppo/PP0_Pendulum-v0_20cbf_00000_0_2021-09-24_15-20-03/\
checkpoint_000029/checkpoint-29
```

Once you figured out the experiment ID and chose a checkpoint ID (as a rule of thumb the larger the ID, the better the results), you can evaluate the training performance of your pendulum training run like this (we'll explain what rollout means in this context in [Chapter 3](#)):

```
rllib rollout \
~/ray_results/pendulum-ppo/PP0_Pendulum-v0_<experiment_id> \
/checkpoint_000<cp-id>/checkpoint-<cp-id> \
--run PP0 --env Pendulum-v0 --steps 2000
```

You should see an animation of a pendulum controlled by an agent that looks like [figure 1-3](#). Since we opted for a quick training procedure instead of maximizing performance, you should see the agent struggle with the pendulum exercise. We could have done much better, and if you're interested to scan Ray's tuned examples for the `Pendulum-v0` environment, you'll find an abundance of solutions to this exercise. The point of this example was to show you how simple it can be to train and evaluate reinforcement learning tasks with RLLib, using just two command line calls to `rllib`.

## Distributed Training with Ray Train

Ray RLLib is dedicated to reinforcement learning, but what do you do if you need to train models for other types of machine learning, like supervised learning? You can use another Ray library for distributed training in this case, called *Ray Train*. At this point, we don't have built up enough knowledge of frameworks such as TensorFlow to give you a concrete and informative example for Ray Train. We'll discuss all of that in [???](#), when it's time to. But we can at least roughly sketch what a distributed training “wrapper” for an ML model would look like, which is simple enough conceptually:

Example 1-6.

```
from ray.train import Trainer
```

```
def training_function(): ❶  
    pass
```

```
trainer = Trainer(backend="tensorflow", num_workers=4) ❷  
trainer.start()
```

```
results = trainer.run(training_function) ❸  
trainer.shutdown()
```

- ❶ First, define your ML model training function. We simply pass here.
- ❷ Then initialize a `Trainer` instance with TensorFlow as the backend.
- ❸ Lastly, scale out your training function on a Ray cluster.

If you're interested in distributed training, you could jump ahead to ???.

## Hyperparameter Tuning

Naming things is hard, but the Ray team hit the spot with *Ray Tune*, which you can use to tune all sorts of parameters. Specifically, it was built to find good hyperparameters for machine learning models. The typical setup is as follows:

- You want to run an extremely computationally expensive training function. In ML it's not uncommon to run training procedures that take days, if not weeks, but let's say you're dealing with just a couple of minutes.
- As result of training, you compute a so-called objective function. Usually you either want to maximize your gains or minimize your losses in terms of performance of your experiment.
- The tricky bit is that your training function might depend on certain parameters, hyperparameters, that influence the value of your objective function.
- You may have a hunch what individual hyperparameters should be, but tuning them all can be difficult. Even if you can restrict these parameters to a sensible range, it's usually prohibitive to test a wide range of combinations. Your training function is simply too expensive.

What can you do to efficiently sample hyperparameters and get “good enough” results on your objective? The field concerned with solving this problem is called *hyperparameter optimization* (HPO), and Ray Tune has an enormous suite of algorithms for

tackling it. Let's look at a first example of Ray Tune used for the situation we just explained. The focus is yet again on Ray and its API, and not on a specific ML task (which we simply simulate for now).

*Example 1-7. Minimizing an objective for an expensive training function with Ray Tune.*

```
from ray import tune
import math
import time

def training_function(config): ❶
    x, y = config["x"], config["y"]
    time.sleep(10)
    score = objective(x, y)
    tune.report(score=score) ❷

def objective(x, y):
    return math.sqrt((x**2 + y**2)/2) ❸

result = tune.run( ❹
    training_function,
    config={
        "x": tune.grid_search([-1, -.5, 0, .5, 1]), ❺
        "y": tune.grid_search([-1, -.5, 0, .5, 1])
    })

print(result.get_best_config(metric="score", mode="min"))
```

- ❶ We simulate an expensive training function that depends on two hyperparameters `x` and `y`, read from a `config`.
- ❷ After sleeping for 5 seconds to simulate training and computing the objective we report back the score to `tune`.
- ❸ The objective computes the mean of the squares of `x` and `y` and returns the square root of this term. This type of objective is fairly common in ML.
- ❹ We then use `tune.run` to initialize hyperparameter optimization on our `training_function`.
- ❺ A key part is to provide a parameter space for `x` and `y` for `tune` to search over.

The Tune example in [Example 1-7](#) finds the best possible choices of parameters `x` and `y` for a `training_function` with a given objective we want to minimize. Even

though the objective function might look a little intimidating at first, since we compute the sum of squares of  $x$  and  $y$ , all values will be non-negative. That means the smallest value is obtained at  $x=0$  and  $y=0$  which evaluates the objective function to 0.

We do a so-called *grid search* over all possible parameter combinations. As we explicitly pass in five possible values for both  $x$  and  $y$  that's a total of 25 combinations that get fed into the training function. Since we instruct `training_function` to sleep for 10 seconds, testing all combinations of hyperparameters sequentially would take more than four minutes total. Since Ray is smart about parallelizing this workload, on my laptop this whole experiment only takes about 35 seconds. Now, imagine each training run would have taken several hours, and we'd have 20 instead of two hyperparameters. That makes grid search infeasible, especially if you don't have educated guesses on the parameter range. In such situations you'll have to use more elaborate HPO methods from Ray Tune, as discussed in ???.

## Model Serving

The last of Ray's high-level libraries we'll discuss specializes on model serving and is simply called *Ray Serve*. To see an example of it in action, you need a trained ML model to serve. Luckily, nowadays you can find many interesting models on the internet that have already been trained for you. For instance, *Hugging Face* has a variety of models available for you to download directly in Python. The model we'll use is a language model called *GPT-2* that takes text as input and produces text to continue or complete the input. For example, you can prompt a question and GPT-2 will try to complete it.

Serving such a model is a good way to make it accessible. You may not now how to load and run a TensorFlow model on your computer, but you do now how to ask a question in plain English. Model serving hides the implementation details of a solution and lets users focus on providing inputs and understanding outputs of a model.

To proceed, make sure to run `pip install transformers` to install the Hugging Face library that has the model we want to use. With that we can now import and start an instance of Ray's serve library, load and deploy a GPT-2 model and ask it for the meaning of life, like so:

*Example 1-8.*

```
from ray import serve
from transformers import pipeline
import requests

serve.start() ❶
```



```

@serve.deployment ❷
def model(request):
    language_model = pipeline("text-generation", model="gpt2") ❸
    query = request.query_params["query"]
    return language_model(query, max_length=100) ❹

model.deploy() ❺

query = "What's the meaning of life?"
response = requests.get(f"http://localhost:8000/model?query={query}") ❻
print(response.text)

```

- ❶ We start serve locally.
- ❷ The `@serve.deployment` decorator turns a function with a `request` parameter into a serve deployment.
- ❸ Loading `language_model` inside the `model` function for every request is inefficient, but it's the quickest way to show you a deployment.
- ❹ We ask the model to give us at most 100 characters to continue our query.
- ❺ Then we formally deploy the model so that it can start receiving requests over HTTP.
- ❻ We use the indispensable `requests` library to get a response for any question you might have.

In ??? you will learn how to properly deploy models in various scenarios, but for now I encourage you to play around with this example and test different queries. Running the last two lines of code repeatedly will give you different answers practically every time. Here's a darkly poetic gem, raising more questions, that I queried on my machine and slightly censored for underaged readers:

```

[[
    "generated_text": "What's the meaning of life?\n\n
    Is there one way or another of living?\n\n
    How does it feel to be trapped in a relationship?\n\n
    How can it be changed before it's too late?\n\n
    What did we call it in our time?\n\n
    Where do we fit within this world and what are we going to live for?\n\n
    My life as a person has been shaped by the love I've received from others."
]]

```

This concludes our whirlwind tour of Ray's data science libraries, the second of Ray's layers. Before we wrap up this chapter, let's have a very brief look at the third layer, the growing ecosystem around Ray.

# A Growing Ecosystem

Ray's high-level libraries are powerful and deserve a much deeper treatment throughout the book. While their usefulness for the data science experimentation lifecycle is undeniable, I also don't want to give off the impression that Ray is all you need from now on. In fact, I believe the best and most successful frameworks are the ones that integrate well with existing solutions and ideas. It's better to focus on your core strengths and leverage other tools for what's missing in your solution. There's usually no reason to re-invent the wheel.

## How Ray Integrates and Extends

To give you an example for how Ray integrates with other tools, consider that Ray Data is a relatively new addition to its libraries. If you want to boil it down, and maybe oversimplify a little, Ray is a compute-first framework. In contrast, distributed frameworks like Apache Spark<sup>7</sup> or Dask can be considered data-first. Pretty much anything you do with Spark starts with the definition of a distributed dataset and transformations thereof. Dask bets on bringing common data structures like Pandas dataframes or Numpy arrays to a distributed setup. Both are immensely powerful in their own regard, and we'll give you a more detailed and fair comparison to Ray in [???](#). The gist of it is that Ray Data does not attempt to replace these tools. Instead, it integrates well with both. As you'll come to see, that's a common theme with Ray.

## Ray as Distributed Interface

One aspect of Ray that's vastly understated in my eyes is that its libraries seamlessly integrate common tools as *backends*. Ray often creates common interfaces, instead of trying to create new standards<sup>8</sup>. These interfaces allow you to run tasks in a distributed fashion, a property most of the respective backends don't have, or not to the same extent. For instance, Ray RLlib and Train are backed by the full power of TensorFlow and PyTorch. Ray Tune supports algorithms from practically every notable HPO tool available, including Hyperopt, Optuna, Nevergrad, Ax, SigOpt and many others. None of these tools are distributed by default, but Tune unifies them in a common interface. Ray Serve can be used with frameworks such as FastAPI, and Ray Data is backed by Arrow and comes with many integrations to other frameworks,

---

<sup>7</sup> Spark has been created by another lab in Berkeley, AMPLab. The internet is full of blog posts claiming that Ray should therefore be seen as a replacement of Spark. It's better to think of them as tools with different strengths that are both likely here to stay.

<sup>8</sup> Before the deep learning framework [Keras](#) became an official part of a corporate flagship, it started out as a convenient API specification for various lower-level frameworks such as Theano, CNTK, or TensorFlow. In that sense Ray RLlib has the chance to become Keras for RL. Ray Tune might just be Keras for HPO. The missing piece for more adoption is probably a more elegant API for both.


such as Spark and Dask. Overall this seems to be a robust design pattern that can be used to extend current Ray projects or integrate new backends in the future.

## Summary

To sum up what we've discussed in this chapter, [Figure 1-4](#) gives you an overview of the three layers of Ray as we laid them out. Ray's core distributed execution engine sits at the center of the framework. For practical data science workflows you can use Ray Data for data processing, Ray RLlib for reinforcement learning, Ray Train for distributed model training, Ray Tune for hyperparameter tuning and Ray Serve for model serving. You've seen examples for each of these libraries and have an idea of what their APIs entail. On top of that, Ray's ecosystem has many extensions that we'll look more into later on. Maybe you can already spot a few tools you know and like in [Figure 1-4](#)<sup>9</sup>?

---

<sup>9</sup> Note that “Ray Train” has been called “raysgd” in older versions of Ray, and does not have a new logo yet.



---

# Getting Started With Ray Core

## A Note for Early Release Readers

With Early Release ebooks, you get books in their earliest form—the author’s raw and unedited content as they write—so you can take advantage of these technologies long before the official release of these titles.

For a book on distributed Python, it’s not without a certain irony that Python on its own is largely ineffective for distributed computing. Its interpreter is effectively single threaded which makes it difficult to, for example, leverage multiple CPUs on the same machine, let alone a whole cluster of machines, using plain Python. That means you need extra tooling, and luckily the Python ecosystem has some options for you. For instance, libraries like `multiprocessing` can help you distribute work on a single machine, but not beyond.

In this chapter you’ll understand how Ray core handles distributed computing by spinning up a local cluster, and you’ll learn how to use Ray’s lean and powerful API to parallelize some interesting computations. For instance, you’ll build an example that runs a data-parallel task efficiently and asynchronously on Ray, in a convenient way that’s not easily replicable with other tooling. We discuss how *tasks* and *actors* work as distributed versions of functions and classes in Python. You’ll also learn about all the fundamental concepts underlying Ray and what its architecture looks like. In other words, we’ll give you a look under the hood of Ray’s engine.

## An Introduction To Ray Core

The bulk of this chapter is an extended Ray core example that we’ll build together. Many of Ray’s concepts can be explained with a good example, so that’s exactly what

we'll do. As before, you can follow this example by typing the code yourself (which is highly recommended), or by following the [notebook for this chapter](#).

In [Chapter 1](#) we've introduced you to the very basics of Ray clusters and showed you how start a local cluster simply by typing

*Example 2-1.*

```
import ray
ray.init()
```

You'll need a running Ray cluster to run the examples in this chapter, so make sure you've started one before continuing. The goal of this section is to give you a quick introduction to the Ray Core API, which we'll simply refer to as the Ray API from now on.

As a Python programmer, the great thing about the Ray API is that it hits so close to home. It uses familiar concepts such as decorators, functions and classes to provide you with a fast learning experience. The Ray API aims to provide a universal programming interface for distributed computing. That's certainly no easy feat, but I think Ray succeeds in this respect, as it provides you with good abstractions that are intuitive to learn and use. Ray's engine does all the heavy lifting for you in the background. This design philosophy is what enables Ray to be used with existing Python libraries and systems.

## A First Example Using the Ray API

To give you an example, take the following function which retrieves and processes data from a database. Our dummy database is a plain Python list containing the words of the title of this book. We act as if retrieving an individual item from this database and further processing it is expensive by letting Python sleep.

*Example 2-2.*

```
import time

database = [ ❶
    "Learning", "Ray",
    "Flexible", "Distributed", "Python", "for", "Data", "Science"
]

def retrieve(item):
    time.sleep(item / 10.) ❷
    return item, database[item]
```

- ❶ A dummy database containing string data with the title of this book.
- ❷ We emulate a data-crunching operation that takes a long time.

Our database has eight items, from `database[0]` for “Learning” to `database[7]` for “Science”. If we were to retrieve all items sequentially, how long should that take? For the item with index 5 we wait for half a second ( $5 / 10.$ ) and so on. In total, we can expect a runtime of around  $(0+1+2+3+4+5+6+7)/10. = 2.8$  seconds. Let’s see if that’s what we actually get:

*Example 2-3.*

```
def print_runtime(input_data, start_time, decimals=1):
    print(f'Runtime: {time.time() - start_time:.{decimals}f} seconds, data:')
    print(*input_data, sep="\n")

start = time.time()
data = [retrieve(item) for item in range(8)] ❶
print_runtime(data, start) ❷
```

- ❶ We use a list comprehension to retrieve all eight items.
- ❷ Then we unpack the data to print each item on its own line.

If you run this code, you should see the following output:

```
Runtime: 2.8 seconds, data:
(0, 'Learning')
(1, 'Ray')
(2, 'Flexible')
(3, 'Distributed')
(4, 'Python')
(5, 'for')
(6, 'Data')
(7, 'Science')
```

We cut off the output of the program after one decimal number. There’s a little overhead that brings the total closer to 2.82 seconds. On your end this might be slightly less, or much more, depending on your computer. The important take-away is that our naive Python implementation is not able to run this function in parallel. This may not come as a surprise to you, but you could have at least suspected that Python list comprehensions are more efficient in that regard. The runtime we got is pretty much the worst case scenario, namely the 2.8 seconds we calculated prior to running the code. If you think about it, it might even be a bit frustrating to see that a program that essentially sleeps most of its runtime is that slow overall. Ultimately you can blame the *Global Interpreter Lock* (GIL) for that, but it gets enough of it already.



## Python's Global Interpreter Lock

The Global Interpreter Lock or GIL<sup>1</sup> is undoubtedly one of the most infamous features of the Python language. In a nutshell it's a lock that makes sure only one thread on your computer can ever execute your Python code at a time. If you use multi-threading, the threads need to take turns controlling the Python interpreter.

The GIL has been implemented for good reasons. For one, it makes memory management that much easier in Python. Another key advantage is that it makes single-threaded programs quite fast. Programs that primarily use lots of system input and output (we say they are I/O-bound), like reading files or databases, benefit as well. One of the major downsides is that CPU-bound programs are essentially single-threaded. In fact, CPU-bound tasks might even run *faster* when not using multi-threading, as the latter incurs write-lock overheads on top of the GIL.

Given all that, the GIL might somewhat paradoxically be one of the reasons for Python's popularity, if you believe [Larry Hastings](#). Interestingly, Hastings also led (unsuccessful) efforts to remove it in a project called *GILectomy*, which is exactly the kind of complicated surgery that it sounds like. The jury is still out, but [Sam Gross](#) might just have found a way to remove the GIL in his `nogil` branch of Python 3.9. For now, if you absolutely have to work around the GIL, consider using an implementation different from CPython. CPython is Python's standard implementation, and if you don't know that you're using it, you're definitely using it. Implementations like Jython, IronPython or PyPy don't have a GIL, but come with their own drawbacks.

## Functions and Remote Ray Tasks

It's reasonable to assume that such a task can benefit from parallelization. Perfectly distributed, the runtime should not take much longer than the longest subtask, namely  $7/10 = 0.7$  seconds. So, let's see how you can extend this example to run on Ray. To do so, you start by using the `@ray.remote` decorator as follows:

*Example 2-4.*

```
@ray.remote ❶
def retrieve_task(item):
    return retrieve(item) ❷
```

---

<sup>1</sup> I still don't know how to pronounce this acronym, but I get the feeling that the same people who pronounce GIF like "giraffe" also say GIL like "guitar". Just pick one, or spell it out, if you feel insecure.



- ❶ With just this decorator we make any Python function a Ray task.
- ❷ All else remains unchanged. `retrieve_task` just passes through to `retrieve`.

In this way, the function `retrieve_task` becomes a so-called Ray task. That's an extremely convenient design choice, as you can focus on your Python code first, and don't have to completely change your mindset or programming paradigm to use Ray. Note that in practice you would have simply added the `@ray.remote` decorator to your original `retrieve` function (after all, that's the intended use of decorators), but we didn't want to touch previous code to keep things as clear as possible.

Easy enough, so what do you have to change in the code that retrieves the data and measures performance? It turns out, not much. Let's have a look at how you'd do that:

*Example 2-5. Measuring performance of your Ray task.*

```
start = time.time()
data_references = [retrieve_task.remote(item) for item in range(8)] ❶
data = ray.get(data_references) ❷
print_runtime(data, start, 2)
```

- ❶ To run `retrieve_task` on your local Ray cluster, you use `.remote()` and pass in your data as before. You'll get a list of object references.
- ❷ To get back data, and not just Ray object references, you use `ray.get`.

Did you spot the differences? You have to execute your Ray task remotely using the `remote` function. When tasks get executed remotely, even on your local cluster, Ray does so *asynchronously*. The list items in `data_references` in the last code snippet do not contain the results directly. In fact, if you check the Python type of the first item with `type(data_references[0])` you'll see that it's in fact an `ObjectRef`. These object references correspond to *futures* which you need to ask the result of. This is what the call to `ray.get(...)` is for.

We still want to work more on this example<sup>2</sup>, but let's take a step back here and recap what we did so far. You started with a Python function and decorated it with `@ray.remote`. This made your function a Ray task. Then, instead of calling the original function in your code straight-up, you called `.remote(...)` on the Ray task. The last step was to `.get(...)` the results back from your Ray cluster. I think this procedure is so intuitive that I'd bet you could already create your own Ray task from

---

<sup>2</sup> This example has been adapted from Dean Wampler's fantastic report "[What is Ray?](#)".

another function without having to look back at this example. Why don't you give it a try right now?

Coming back to our example, by using Ray tasks, what did we gain in terms of performance? On my machine the runtime clocks in at 0.71 seconds, which is just slightly more than the longest subtask, which comes in at 0.7 seconds. That's great and much better than before, but we can further improve our program by leveraging more of Ray's API.

## Using the object store with put and get

One thing you might have noticed is that in the definition of `retrieve` we *directly* accessed items from our database. Working on a local Ray cluster this is fine, but imagine you're running on an actual cluster comprising several computers. How would all those computers access the same data? Remember from [Chapter 1](#) that in a Ray cluster there is one head node with a driver process (running `ray.init()`) and many worker nodes with worker processes executing your tasks. My laptop has a total of 8 CPU cores, so Ray will create 8 worker processes on my one-node local cluster. Our database is currently defined on the driver only, but the workers running your tasks need to have access to it to run the `retrieve` task. Luckily, Ray provides an easy way to share data between the driver and workers (or between workers). You can simply use `put` to place your data into Ray's *distributed object store* and then use `get` on the workers to retrieve it as follows.

*Example 2-6.*

```
database_object_ref = ray.put(database) ❶
```

```
@ray.remote
def retrieve_task(item):
    obj_store_data = ray.get(database_object_ref) ❷
    time.sleep(item / 10.)
    return item, obj_store_data[item]
```

- ❶ put your database into the object store and receive a reference to it.
- ❷ This allows your workers to get the data, no matter where they are located in the cluster.

By using the object store this way, you can let Ray handle data access across the whole cluster. We'll talk about how exactly data is passed between nodes and within workers when talking about Ray's infrastructure. While the interaction with the object store requires some overhead, Ray is really smart about storing the data, which gives you performance gains when working with larger, more realistic datasets. For now, the

important part is that this step is essential in a truly distributed setting. If you like, try to re-run [Example 2-5](#) with this new `retrieve_task` function and confirm that it still runs, as expected.

## Using Ray's `wait` function for non-blocking calls

Note how in [Example 2-5](#) we used `ray.get(data_references)` to access results. This call is *blocking*, which means that our driver has to wait for all the results to be available. That's not a big deal in our case, the program now finishes in under a second. But imagine processing of each data item would take several minutes. In that case you would want to free up the driver process for other tasks, instead of sitting idly by. Also, it would be great to process results as they come in (some finish much quicker than others), rather than waiting for all data to be processed. One more question to keep in mind is what happens if one of the data items can't be retrieved as expected? Let's say there's a deadlock somewhere in the database connection. In that case, the driver will simply hang and never retrieve all items. For that reason it's a good idea to work with reasonable timeouts. In our scenario, we should not wait longer than 10 times the longest data retrieval task before stopping the task. Here's how you can do that with Ray by using `wait`:

*Example 2-7.*

```
start = time.time()
data_references = [retrieve_task.remote(item) for item in range(8)]
all_data = []

while len(data_references) > 0: ❶
    finished, data_references = ray.wait(data_references, num_returns=2, timeout=7.0) ❷
    data = ray.get(finished)
    print_runtime(data, start, 3) ❸
    all_data.extend(data) ❹
```

- ❶ Instead of blocking, we loop through unfinished `data_references`.
- ❷ We asynchronously wait for finished data with a reasonable timeout. `data_references` gets overridden here, to prevent an infinite loop.
- ❸ We print results as they come in, namely in blocks of two.
- ❹ Then we append new data to the `all_data` until finished.

As you can see `ray.wait` returns two arguments, namely finished data and futures that still need to be processed. We use the `num_returns` argument, which defaults to 1, to let `wait` return whenever a new pair of data items is available. On my laptop this results in the following output:

```

Runtime: 0.108 seconds, data:
(0, 'Learning')
(1, 'Ray')
Runtime: 0.308 seconds, data:
(2, 'Flexible')
(3, 'Distributed')
Runtime: 0.508 seconds, data:
(4, 'Python')
(5, 'for')
Runtime: 0.709 seconds, data:
(6, 'Data')
(7, 'Science')

```

Note how in the while loop, instead of just printing results, we could have done many other things, like starting entirely new tasks on other workers with the data already retrieved up to this point.

## Handling Task Dependencies

So far our example program has been fairly easy on a conceptual level. It consists of a single step, namely retrieving a bunch of data. Now, imagine that once your data is loaded you want to run a follow-up processing task on it. To be more concrete, let's say we want to use the result of our first retrieve task to query other, related data (pretend that you're querying data from a different table in the same database). The following code sets up such a task and runs both our `retrieve_task` and `follow_up_task` consecutively.

*Example 2-8. Running a follow-up task that depends on another Ray task.*

```

@ray.remote
def follow_up_task(retrieve_result): ❶
    original_item, _ = retrieve_result
    follow_up_result = retrieve(original_item + 1) ❷
    return retrieve_result, follow_up_result ❸

retrieve_refs = [retrieve_task.remote(item) for item in [0, 2, 4, 6]]
follow_up_refs = [follow_up_task.remote(ref) for ref in retrieve_refs] ❹

result = [print(data) for data in ray.get(follow_up_refs)]

```

- ❶ Using the result of `retrieve_task` we compute another Ray task on top of it.
- ❷ Leveraging the `original_item` from the first task, we retrieve more data.
- ❸ Then we return both the original and the follow-up data.
- ❹ We pass the object references from the first task into the second task.

Running this code results in the following output.

```
((0, 'Learning'), (1, 'Ray'))  
((2, 'Flexible'), (3, 'Distributed'))  
((4, 'Python'), (5, 'for'))  
((6, 'Data'), (7, 'Science'))
```

If you don't have a lot of experience with asynchronous programming, you might not be impressed by [Example 2-8](#). But I hope to convince you that it's at least a bit surprising<sup>3</sup> that this code snippet runs at all. So, what's the big deal? After all, the code reads like regular Python - a function definition and a few list comprehensions. The point is that the function body of `follow_up_task` expects a Python tuple for its input argument `retrieve_result`, which we unpack in the first line of the function definition.

But by invoking `[follow_up_task.remote(ref) for ref in retrieve_refs]` we do *not* pass in tuples to the follow-up task at all. Instead, we pass in Ray *object references* with `retrieve_refs`. What happens under the hood is that Ray knows that `follow_up_task` requires actual values, so internally in this task it will call `ray.get` to resolve the futures. Ray builds a dependency graph for all tasks and executes them in an order that respects the dependencies. You do not have to tell Ray explicitly when to wait for a previous task to finish, it will infer that information for you.

The follow-up tasks will only be scheduled, once the individual retrieve tasks have finished. If you ask me, that's an incredible feature. In fact, if I had called `retrieve_refs` something like `retrieve_result`, you may not even have noticed this important detail. That's by design. Ray wants you to focus on your work, not on the details of cluster computing. In [figure 2-1](#) you can see the dependency graph for the two tasks visualized.

---

<sup>3</sup> According to [Clarke's third law](#) any sufficiently advanced technology is indistinguishable from magic. For me, this example has a bit of magic to it.

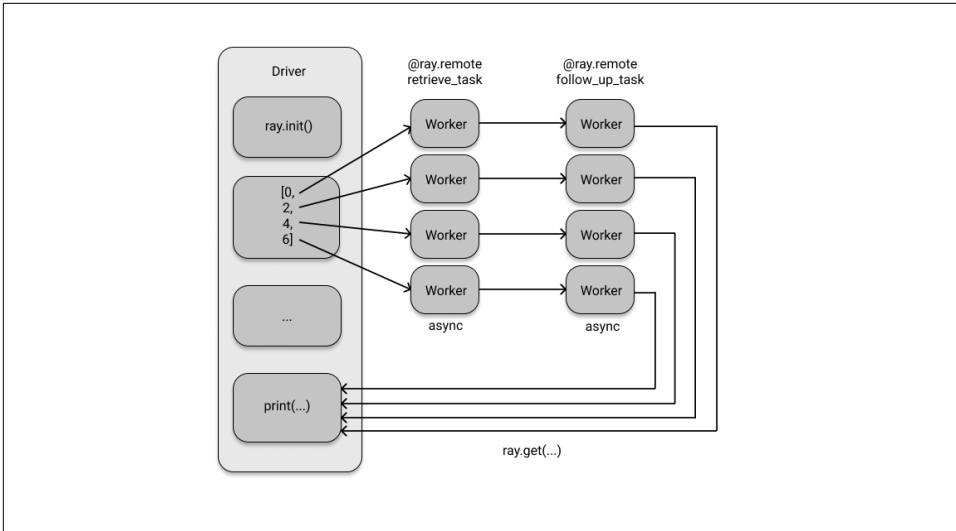


Figure 2-1. Running two dependent tasks asynchronously and in parallel with Ray.

If you feel like it, try to rewrite [Example 2-8](#) so that it explicitly uses `get` on the first task before passing values into the follow-up task. That does not only introduce more boilerplate code, but it's also a bit less intuitive to write and understand.

## From Classes to Actors

Before wrapping up this example, let's discuss one more important concept of Ray Core. Notice how in our example everything is essentially a function. We just used the `ray.remote` decorator to make some of them remote functions, and other than that simply used plain Python. Let's say we wanted to track how often our database has been queried? Sure, we could simply count the results of our retrieve tasks, but is there a better way to do this? We want to track this in a “distributed” way that will scale. For that, Ray has the concept of *actors*. Actors allow you to run *stateful* computations on your cluster. They can also communicate between each other<sup>4</sup>. Much like Ray tasks were simply decorated functions, Ray actors are decorated Python classes. Let's write a simple counter to track our database calls.

*Example 2-9.*

```
@ray.remote ❶
class DataTracker:
    def __init__(self):
```

<sup>4</sup> The actor model is an established concept in computer science, which you can find implemented e.g. in Akka or Erlang. However, the history and specifics of actors are not relevant to our discussion.

```

self._counts = 0

def increment(self):
    self._counts += 1

def counts(self):
    return self._counts

```

- ❶ We can make any Python class a Ray actor by using the same `ray.remote` decorator as before.

This `DataTracker` class is already an actor, since we equipped it with the `ray.remote` decorator. This actor can track state, here just a simple counter, and its methods are Ray tasks that get invoked precisely like we did with functions before, namely using `.remote()`. Let's see how we can modify our existing `retrieve_task` to incorporate this new actor.

*Example 2-10.*

```

@ray.remote
def retrieve_tracker_task(item, tracker): ❶
    obj_store_data = ray.get(database_object_ref)
    time.sleep(item / 10.)
    tracker.increment.remote() ❷
    return item, obj_store_data[item]

tracker = DataTracker.remote() ❸

data_references = [retrieve_tracker_task.remote(item, tracker) for item in range(8)] ❹
data = ray.get(data_references)
print(ray.get(tracker.counts.remote())) ❺

```

- ❶ We pass in the tracker actor into this task.
- ❷ The tracker receives an increment for each call.
- ❸ We instantiate our `DataTracker` actor by calling `.remote()` on the class.
- ❹ The actor gets passed into the retrieve task.
- ❺ Afterwards we can get the counts state from our tracker from another remote invocation.

Unsurprisingly, the result of this computation is in fact 8. We didn't need actors to compute this, but I hope you can see how useful it can be to have a mechanism to track state across the cluster, potentially spanning multiple tasks. In fact, we could

pass our actor into any dependent task, or even pass it into the constructor of yet another actor. There is no limitation to what you can do, and it's this flexibility that makes the Ray API so powerful. It's also worth mentioning that it's not very common for distributed Python tools to allow for stateful computations like this. This feature can come in very handy, especially when running complex distributed algorithms, for instance when using reinforcement learning. This completes our extensive first Ray API example. Let's see if we can concisely summarize the Ray API next.

## An Overview of the Ray Core API

If you recall what exactly we did in the previous example, you'll notice that we used a total of just six API methods<sup>5</sup>. You used `ray.init()` to start the cluster and `@ray.remote` to turn functions and classes into tasks and actors. Then we used `ray.put()` to pass data into Ray's object store and `ray.get()` to retrieve data from the cluster. Finally, we used `.remote()` on actor methods or tasks to run code on our cluster, and `ray.wait` to avoid blocking calls.

While six API methods might not seem like much, those are the only ones you'll likely ever care about when using the Ray API<sup>6</sup>. Let's briefly summarize them in a table, so you can easily reference them in the future.

Table 2-1. The six major API methods of Ray Core.

API call	Description
<code>ray.init()</code>	Initializes your Ray cluster. Pass in an address to connect to an existing cluster.
<code>@ray.remote</code>	Turns functions into tasks and classes into actors.
<code>ray.put()</code>	Puts data into Ray's object store.
<code>ray.get()</code>	Gets data from the object store. Returns data you've put there or that was computed by a task or actor.
<code>.remote()</code>	Runs actor methods or tasks on your Ray cluster and is used to instantiate actors.
<code>ray.wait()</code>	Returns two lists of object references, one with finished tasks we're waiting for and one with unfinished tasks.

Now that you've seen the Ray API in action, let's quickly discuss Ray's design philosophy, before moving on to discussing its system architecture.

<sup>5</sup> To paraphrase [Alan Kay](#), to get simplicity, you need to find slightly more sophisticated building blocks. In my eyes, the Ray API does just that for distributed Python.

<sup>6</sup> You can check out the [API reference](#) to see that there are in fact quite a bit more methods available. At some point you should invest in understanding the arguments of `init`, but all other methods likely won't be of interest to you, if you're not an administrator of your Ray cluster.



## Design Principles

Ray is built with several design principles in mind, most of which you’ve got a taste of already. Its API is designed for simplicity and generality. Its compute model banks on flexibility. And its system architecture is designed for performance and scalability. Let’s look at each of these in more detail.

### Simplicity & Abstraction

As you’ve seen, Ray’s API does not only bank on simplicity, it’s also intuitive to pick up. It doesn’t matter whether you just want to use all the CPU cores on your laptop or leverage all the machines in your cluster. You might have to change a line of code or two, but the Ray code you use stays essentially the same. And as with any good distributed system, Ray manages task distribution and coordination under the hood. That’s great, because you’re not bogged down by reasoning about the mechanics of distributed computing. A good abstraction layer allows you to focus on your work, and I think Ray has done a great job of giving you one.

Since Ray’s API is so generally applicable and pythonic, it’s easy to integrate with other tools. For instance, Ray actors can call into or be called by existing distributed Python workloads. In that sense Ray makes for good “glue code” for distributed workloads, too, as its performant and flexible enough to communicate between different systems and frameworks.

### Flexibility

For AI workloads, in particular when dealing with paradigms like reinforcement learning, you need a flexible programming model. Ray’s API is designed to make it easy to write flexible and composable code. Simply put, if you can express your workload in Python, you can distribute it with Ray. Of course, you still need to make sure you have enough resources available and be mindful of what you want to distribute. But Ray doesn’t limit you in what you can do with it.

Ray is also flexible when it comes to *heterogeneity* of computations. For instance, let’s say you work on a complex simulation. Simulations can usually be decomposed into several tasks or steps. Some of these steps might take hours to run, others just a few milliseconds, but they always need to be scheduled and executed quickly. Sometimes a single task in a simulation can take a long time, but other, smaller tasks should be able to run in parallel without blocking it. Also, subsequent tasks may depend on the outcome of an upstream task, so you need a framework to allow for *dynamic execution* that deals well with task dependencies. In the example we discussed in this chapter you’ve seen that Ray’s API is built for that.

You also need to ensure you are flexible in your resource usage. For instance, some tasks might have to run on a GPU, while others run best on a couple of CPU cores. Ray provides you with that flexibility.

## Speed & Scalability

Another of Ray's design principles is the speed at which Ray executes its heterogeneous tasks. It can handle millions of tasks per second. What's more is that you only incur very low latencies with Ray. It's built to execute its tasks with just milliseconds of latency.

For a distributed system to be fast, it also needs to scale well. Ray is efficient at distributing and scheduling your tasks across your compute cluster. And it does so in a fault tolerant way, too. In distributed systems it's not a question of if, but when things go wrong. A machine might have an outage, abort a task or simply go up in flames.<sup>7</sup> In any case, Ray is built to recover quickly from failures, which contributes to its overall speed.

# Understanding Ray System Components

You've seen how the Ray API can be used and understand the design philosophy behind Ray. Now it's time to get a better understanding of the underlying system components. In other words, how does Ray work and how does it achieve what it does?

## Scheduling and Executing Work on a Node

You know that Ray clusters consist of nodes. We'll first look at what happens on individual nodes, before we zoom out and discuss how the whole cluster interoperates.

As we've already discussed, a worker node consists of several worker processes or simply workers. Each worker has a unique ID, an IP address and a port by which they can be referenced. Workers are called as they are for a reason, they're components that blindly execute the work you give them. But who tells them what to do and when? A worker might be busy already, it may not have the proper resources to run a task (e.g. access to a GPU), and it might not even have the data it needs to run a given task. On top of that, workers have no knowledge of what happens before or after they've executed their workload, there's no coordination.

---

<sup>7</sup> This might sound drastic, but it's not a joke. To name just one example, in March 2021 a French data center powering millions of websites burnt down completely, which you can read about [in this article](#). If your whole cluster burns down, I'm afraid Ray can't help you.

To address these issues, each worker node has a component called *Raylet*. Think of Raylets as the smart components of a node, which manage the worker processes. Raylets are shared between jobs and consist of two components, namely a task scheduler and an object store.

Let's talk about object stores first. In the running example in this chapter we've already used the concept of an object store loosely, without really specifying it. Each node of a Ray cluster is equipped with an object store, within that node's Raylet, and all object stores collectively form the distributed object store of a cluster. An object store has *shared memory* across the node, so that each worker process has easy access to it. The object store is implemented in **Plasma**, which now belongs to the Apache Arrow project. Functionally, the object store takes care of memory management and ultimately makes sure workers have access to the data they need.

The second component of a Raylet is its scheduler. The scheduler takes care of resource management, among other things. For instance, if a task requires access to 4 CPUs, the scheduler needs to make sure it can find a free worker process that it can *grant access* to said resources. By default, the scheduler knows about and acquires information about the number of CPUs and GPUs and the amount of memory available on its node, but you can register custom resources, if you want to. If it can't provide the required resources, it can't schedule execution of a task.

Apart from resources, the other requirement the scheduler takes care of is *dependency resolution*. That means it needs to ensure that each worker has all the input data it needs to execute a task. For that to work, the scheduler will first resolve local dependencies by looking up data in its object store. If the required data is not available on this node's object store, the scheduler will have to communicate with other nodes (we'll tell you how in a bit) and pull in remote dependencies. Once the scheduler has ensured enough resources for a task, resolved all needed dependencies, and found a worker for a task, it can schedule said task for execution.

Task scheduling is a very difficult topic, even if we're only talking about single nodes. I think you can easily imagine scenarios in which an incorrectly or naively planned task execution can "block" downstream tasks because there are not enough resources left. Especially in a distributed context assigning work like this can become very tricky very quickly.

Now that you know about Raylets, let's briefly come back to worker processes, so that we can wrap up the discussion around worker nodes. An important concept that contributed to the performance of Ray overall is that of *ownership*.

Ownership means that a process that runs something is responsible for it. This makes for a decentralized overall design, since individual tasks have a unique owner. In concrete terms this means that each worker process owns the tasks it submits, which includes proper execution and availability of results (i.e., correct resolution of object

references). Also, anything that gets registered through `ray.put()` is owned by the caller. You should understand ownership in contrast to dependency, which we've already covered by example when discussing task dependencies.

To give you a concrete example, let's say we have a program that starts a task which takes an input value `val` and internally calls another task. That could look as follows:

*Example 2-11.*

```
@ray.remote
def task_owned():
    return

@ray.remote
def task(dependency):
    res_owned = task_owned.remote()
    return

val = ray.put("value")
res = task.remote(dependency=val)
```

From this point on we won't mention it again, but this example assumes that you have a running Ray cluster started with `ray.init()`. Let's quickly analyse ownership and dependency for this example. We defined two tasks in `task` and `task_owned`, and we have three variables in total, namely `val`, `res` and `res_owned`. Our main program defines both `val` (which puts "value" into the object store) and `res`, the final result of the whole program, and it also calls `task`. In other words, the driver *owns* `task`, `val` and `res` according to Ray's ownership definition. In contrast, `res` depends on `task`, but there's no ownership relationship between the two. When `task` gets called, it takes `val` as a dependency. It then calls `task_owned` and assigns `res_owned`, and hence owns them both. Lastly, `task_owned` itself does not own anything, but certainly `res_owned` depends on it.

Ownership is important to know about, but it's not a concept you encounter all that often when working with Ray. The reason we brought it up in this context is that worker processes need to track what they own. In fact, they possess a so-called *ownership table* exactly for that reason. If a task fails and needs to be recomputed, the worker already owns all the information it needs to do so. On top of that, workers also have an in-process store for small objects, which has a default limit of 100KB. Workers have that store so that small data can be directly accessed and stored without incurring communication overhead with the Raylet object store, which is reserved for large objects.

To sum up this discussion about worker nodes, figure [Figure 2-2](#) gives you an overview of all involved components.

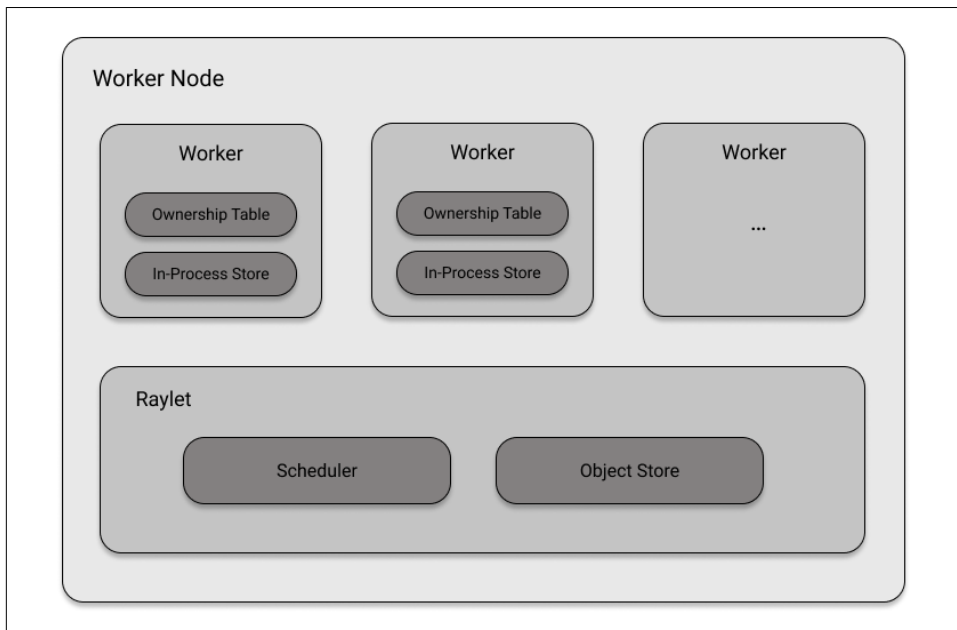


Figure 2-2. The system components comprising a Ray worker node.

## The Head Node

We've already indicated in [Chapter 1](#) that each Ray cluster has one special node called head node. So far you know that this is the node that has the driver process<sup>8</sup>. Drivers can submit tasks themselves, but can't execute them. You also know that the head node can have some worker processes, which is important to be able to run local clusters consisting of a single node. In other words, the head node has everything a worker node has (including a Raylet), but it also has a driver process.

Additionally, the head node comes with a component called *Global Control Store* (GCS). The GCS is a key-value store currently implemented in Redis. It's an important component that carries global information about the cluster, such as system-level metadata. For instance, it has a table with heart beat signals for each Raylet, to ensure they are still reachable. Raylets, in turn, send heart beat signals to the GCS to indicate that they are alive. The GCS also stores the locations of Ray actors and large objects in respective tables, and knows about the dependencies between objects.

<sup>8</sup> In fact, it could have multiple drivers, but this is inessential for our discussion.

## Distributed Scheduling and Execution

Let's briefly talk about cluster orchestration and how nodes manage, plan and execute tasks. When talking about worker nodes, we've indicated that there are several components to distributing workloads with Ray. Here's an overview of the steps and intricacies involved in this process.

- **Distributed memory:** The object stores of individual Raylets share their memory on a node. But sometimes data needs to be transferred between nodes, which is called *distributed object transfer*. This is needed for remote dependency resolution, so that workers have the data they need to run tasks.
- **Communication:** Most of the communication in a Ray cluster, such as object transfer, takes place via the *gRPC* protocol.
- **Resource management and fulfillment:** On a node, Raylets are responsible to grant resources and *lease* worker processes to task owners. All schedulers across nodes form the distributed scheduler. Through communication with the GCS, local schedulers know about other nodes' resources.
- **Task execution:** Once a task has been submitted for execution, all its dependencies (local and remote data) need to be resolved, e.g. by retrieving large data from the object store, before execution can begin.

If the last few sections seem a bit involved technically, that's because they are. In my view it's important to understand the basic patterns and ideas of the software you're using, but I'll admit that the details of Ray's architecture can be a bit tough to wrap your head around in the beginning. In fact, it's one of Ray's design principles to trade-off usability for architectural complexity. If you want to delve deeper into Ray's architecture, a good place to start is [their architecture white paper](#).

To wrap things up, let's summarize all we know in a concise architecture overview with figure [Figure 2-3](#):

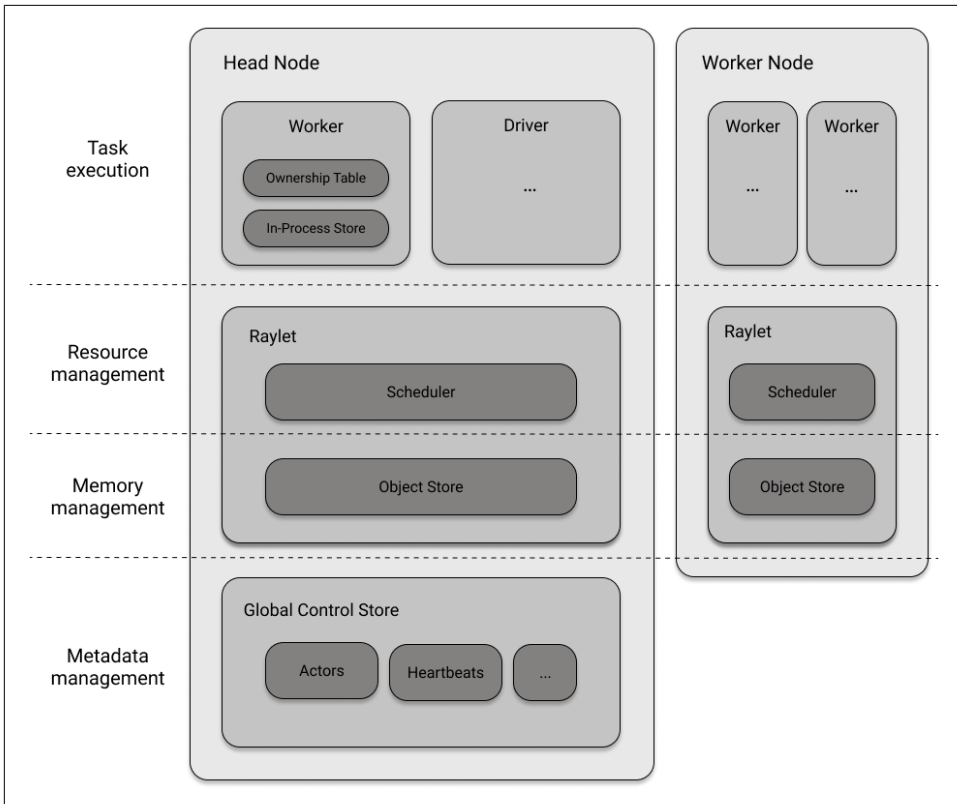


Figure 2-3. An overview of Ray's architectural components.



## Systems related to Ray.

With the architecture and functionality of it in mind, how does Ray relate to other systems? We're not going into the details here, but just touch on the most important topics in broad strokes. Ray can be used as a parallelization framework for Python, and shares properties with tools like `celery` or `multiprocessing`. In fact, there's a **drop-in replacement** for the latter implemented in Ray. Ray is also related to data processing frameworks such as Spark, Dask, Flink or MARS. We'll explore this relationship in ???, when talking about Ray's ecosystem. As a distributed computing tool, Ray also has to deal with the problems of cluster management and orchestration, and we'll see how Ray does that in relation to tools like Kubernetes in ???. Since Ray is implementing the actor model of concurrency, it's also interesting to explore its relationship with frameworks like Akka. Lastly, since Ray banks on a performant, low-level API for communication, there's a certain relationship with high-performance computing (HPC) frameworks and communication protocols like the message passing interface (MPI).

## Summary

You've seen the basics of the Ray API in action in this chapter. You know how to put data to the object store, and how to get it back. Also, you're familiar with declaring Python functions as Ray tasks with the `@ray.remote` decorator, and you know how to run them on a Ray cluster with the `.remote()` call. In much the same way, you understand how to declare a Ray actor from a Python class, how to instantiate it and leverage it for stateful, distributed computations.

On top of that, you also know the basics of Ray clusters. After starting them with `ray.init(...)` you know that you can submit jobs consisting of tasks to your cluster. The driver process, sitting on the head node, will then distribute the tasks to the worker nodes. Raylets on each node will schedule the tasks and worker processes will execute them. This quick tour through Ray core should get you started with writing your own distributed programs, and in the next chapter we'll test your knowledge by implementing a basic machine learning application together.



---

# Building Your First Distributed Application With Ray Core

## A Note for Early Release Readers

With Early Release ebooks, you get books in their earliest form—the author’s raw and unedited content as they write—so you can take advantage of these technologies long before the official release of these titles.

Now that you’ve seen the basics of the Ray API in action, let’s build something more realistic with it. By the end of this comparatively short chapter, you will have built a reinforcement learning problem from scratch, implemented your first algorithm to tackle it, and used Ray tasks and actors to parallelize this solution to a local cluster—all in less than 250 lines of code.

This chapter is designed to work for readers who don’t have any experience with reinforcement learning. We’ll work on a straightforward problem and develop the necessary skills to tackle it hands-on. Since chapter ??? is devoted entirely to this topic, we’ll skip all advanced RL topics and language and just focus on the problem at hand. But even if you’re a quite advanced RL user, you’ll likely benefit from implementing a classical algorithm in a distributed setting.

This is the last chapter working *only* with Ray Core. I hope you learn to appreciate how powerful and flexible it is, and how quickly you can implement distributed experiments, that would otherwise take considerable efforts to scale.

# Creating A Ray Core App

As with the chapters before, I encourage you to code this chapter with me and build this application together as we go. In case you don't want to do that, you can also simply follow [the notebook for this chapter](#).

To give you an idea, the app we're building is structured as follows:

- You implement a simple 2D-maze game in which a single player can move around in the four major directions.
- You initialize the maze as a 5x5 grid to which the player is confined.
- One of the 25 grid cells is the “goal” that a player called the “seeker” must reach.
- Instead of hard-coding a solution, you will employ a reinforcement learning algorithm, so that the seeker learns to find the goal.
- This is done by repeatedly running simulations of the maze, rewarding the seeker for finding the goal and smartly keeping track of which decisions of the seeker worked and which didn't.
- As running simulations can be parallelized and our RL algorithm can also be trained in parallel, we utilize the Ray API to parallelize the whole process.

We're not quite ready to deploy this application on an actual Ray cluster comprised of multiple nodes just yet, so for now we'll continue to work with local clusters. If you're interested in infrastructure topics and want to learn how to set up Ray clusters, jump ahead to [???](#), and to see a fully deployed Ray application you can go to [???](#).

## A Simple Maze Problem

Let's start by implementing the 2D maze we just sketched. The idea is to implement a simple grid in Python that spans a 5x5 grid starting at (0, 0) and ending at (4, 4) and properly define how a player can move around the grid. To do this, we first need an abstraction for moving in the four cardinal directions. These four actions, namely moving up, down, left, and right, can be encoded in Python as a class we call `Discrete`. In the example we're building we only have four discrete actions, but I hope you forgive me the straightforward generalization to an arbitrary number of directions (which we'll actually need):

*Example 3-1.*

```
import random
```

```
class Discrete:
    def __init__(self, num_actions: int):
```

```

        """ Discrete action space for num_actions.
        Discrete(4) can be used as encoding moving in one of the cardinal directions.
        """
        self.n = num_actions

    def sample(self):
        return random.randint(0, self.n - 1) ❶

space = Discrete(4)
print(space.sample()) ❷

```

- ❶ A discrete action can be uniformly sampled between 0 and n-1.
- ❷ For instance, a `Discrete(4)` sample will give you 0, 1, 2, or 3.

Sampling from a `Discrete(4)` like in this example will randomly return 0, 1, 2, or 3. How we interpret these numbers is up to us, so let's say we go for “down”, “left”, “right”, and “up” in that order.

Now that we know how to encode moving around the maze, let's code the maze itself, including the goal cell and the position of the seeker player that tries to find the goal. To this end we're going to implement a Python class called `Environment`. It's called like that, because the maze is the environment in which the player “lives”. To make matters easy, we'll always put the seeker at (0, 0) and the goal at (4, 4). To make the seeker move and find its goal, we initialize the `Environment` with an `action_space` of `Discrete(4)`.

There is one last bit of information we need to set up for our maze environment, namely an encoding of the seeker position. The reason for that is that we're going to implement an algorithm later that keeps track of which actions led to good results for which seeker positions. By encoding the seeker position as a `Discrete(5*5)`, it becomes a single number that's much easier to work with. In RL lingo it is common to call the information of the game accessible to the player an *observation*. So, in analogy to the actions we can carry out for our seeker, we can also define an `observation_space` for it. Here's the implementation of what we've just discussed:

*Example 3-2.*

```

import os

class Environment:

    seeker, goal = (0, 0), (4, 4) ❶
    info = {'seeker': seeker, 'goal': goal}

```

```
def __init__(self, *args, **kwargs):
    self.action_space = Discrete(4) ❷
    self.observation_space = Discrete(5*5) ❸
```

- ❶ The seeker gets initialized in the top left, the goal in the bottom right of the maze.
- ❷ Our seeker agent can move down, left, up and right.
- ❸ And it can be in a total of 25 states, one for each position on the grid.

Note that we defined an `info` variable as well, which can be used to print information about the current state of the maze, for instance for debugging purposes. To play an actual game of find-the-goal from the perspective of the seeker, we have to define a few helper methods. Clearly, the game should be considered “done” when the seeker finds the goal. Also, we should reward the seeker for finding the goal. And when the game is over, we should be able to reset it to its initial state, to play again. To round things off, we also define a `get_observation` method that returns the encoded seeker position. Continuing our implementation of the `Environment` class, this translates into the following four methods.

*Example 3-3.*

```
def reset(self): ❶
    """Reset seeker and goal positions, return observations."""
    self.seeker = (0, 0)
    self.goal = (4, 4)

    return self.get_observation()

def get_observation(self):
    """Encode the seeker position as integer"""
    return 5 * self.seeker[0] + self.seeker[1] ❷

def get_reward(self):
    """Reward finding the goal"""
    return 1 if self.seeker == self.goal else 0 ❸

def is_done(self):
    """We're done if we found the goal"""
    return self.seeker == self.goal ❹
```

- ❶ To play a new game, we’ll have to reset the grid to its original state.
- ❷ Converting the seeker tuple to a value from the environment’s `observation_space`.

- ③ The seeker is only rewarded when reaching the goal.
- ④ If the seeker is at the goal, the game is over.

The last essential method to implement is the `step` method. Imagine you're playing our maze game and decide to go right as your next move. The `step` method will take this action (namely 3, the encoding of “right”) and apply it to the internal state of the game. To reflect what changed, the `step` method will then return the seeker's observations, its reward, whether the game is over, and the `info` value of the game. Here's how the `step` method works:

*Example 3-4.*

```
def step(self, action):
    """Take a step in a direction and return all available information."""
    if action == 0: # move down
        self.seeker = (min(self.seeker[0] + 1, 4), self.seeker[1])
    elif action == 1: # move left
        self.seeker = (self.seeker[0], max(self.seeker[1] - 1, 0))
    elif action == 2: # move up
        self.seeker = (max(self.seeker[0] - 1, 0), self.seeker[1])
    elif action == 3: # move right
        self.seeker = (self.seeker[0], min(self.seeker[1] + 1, 4))
    else:
        raise ValueError("Invalid action")

    return self.get_observation(), self.get_reward(), self.is_done(), self.info ①
```

- ① After taking a step in the specified direction, we return observation, reward, whether we're done, and any additional information we might find useful.

I said the `step` method was the last essential method, but we actually want to define one more helper method that's extremely helpful to visualize the game and help us understand it. This `render` method will print the current state of the game to the command line.

*Example 3-5.*

```
def render(self, *args, **kwargs):
    """Render the environment, e.g. by printing its representation."""
    os.system('cls' if os.name == 'nt' else 'clear') ①
    grid = [['| ' for _ in range(5)] + ["|\n"] for _ in range(5)]
    grid[self.goal[0]][self.goal[1]] = '|G'
    grid[self.seeker[0]][self.seeker[1]] = '|S' ②
    print(''.join([''.join(grid_row) for grid_row in grid])) ③
```

- ❶ First we clear the screen.
- ❷ Then we draw the grid and mark the goal as G and the seeker as S on it.
- ❸ The grid then gets rendered by printing it to your screen.

Great, now we have completed the implementation of our `Environment` class that's defining our 2D-maze game. We can step through this game, know when it's done and reset it again. The player of the game, the seeker, can also observe its environment and gets rewarded for finding the goal.

Let's use this implementation to play a game of find-the-goal for a seeker that simply takes random actions. This can be done by creating a new `Environment`, sampling and applying actions to it, and rendering the environment until the game is over:

*Example 3-6.*

```
import time

environment = Environment()

while not environment.is_done():
    random_action = environment.action_space.sample() ❶
    environment.step(random_action)
    time.sleep(0.1)
    environment.render() ❷
```

- ❶ We can test our environment by applying sampled actions until we're done.
- ❷ To visualize the environment we render it after waiting for a tenth of a second (otherwise the code runs too fast to follow).

If you run this on your computer, eventually you'll see that the game is over and the seeker has found the goal. It might take a while if you're unlucky.

In case you're objecting that this is an extremely simple problem, and to solve it all you have to do is go right and down four times each, I'm not arguing with you. The point is that we want to tackle this problem using machine learning. Specifically, we want to implement an algorithm that figures out on its own how to play the game, merely by playing the game repeatedly: observing what's happening, deciding what to do next, and getting rewarded for your actions.

If you want to now is a good time to make the game more complex yourself. As long as you do not change the interface we defined for the `Environment` class, you could modify this game in many ways. Here are a few suggestions:

- Make it a 10x10 grid or randomize the initial position of the seeker.
- Make the outer walls of the grid dangerous. Whenever you try to touch them, you'll incur a reward of -100, i.e a steep penalty.
- Introduce obstacles in the grid that the seeker cannot pass through.

If you're feeling really adventurous, you could also randomize the goal position. This requires extra care, as currently the seeker has no information about the goal position in terms of the `get_observation` method. Maybe come back to tackling this last exercise after you've finished reading this chapter.

## Building a Simulation

With the `Environment` class implemented, what does it take to tackle the problem of “teaching” the seeker to play the game well and find the goal consistently in the minimum number of 8 steps necessary? We've equipped the maze environment with reward information, so that the seeker can use this signal to learn to play the game. In reinforcement learning, you play games repeatedly and learn from the experience you made in the process. The player of the game is often referred to as *agent* that takes *actions* in the environment, observes its *state* and receives a `_reward_`<sup>1</sup>. The better an agent learns, the better it becomes at interpreting the current game state (observations) and finding actions that lead to more rewarding outcomes.

Regardless of the RL algorithm you want to use (in case you know any), you need to have a way of simulating the game repeatedly, to collect experience data. For this reason we're going to implement a simple `Simulation` class in just a bit.

The other useful abstraction we need to proceed is that of a `Policy`, a way of specifying actions. Right now the only thing we can do to play the game is sampling random actions for our seeker. What a `Policy` allows us to do is to get better actions for the current state of the game. In fact, we define a `Policy` to be a class with a `get_action` method that takes a game state and returns an action.

Remember that our game has a total of 25 seeker states, and 4 actions. A simple idea would be to look at pairs of states and actions and assign a high value to a pair if carrying out this action in this state will lead to a high reward, and a low value otherwise. For instance, from your intuition of the game it should be clear that going down or right is always a good idea, whereas going left or up is not. Then, create a 25x4 lookup table of all possible state-action pairs and store it in our `Policy`. Then we could simply ask our policy to return the highest value of any action given a state. Of

---

<sup>1</sup> As we'll see in chapter ???, you can run RL on multi-player games, too. Making the maze environment a so-called multi-agent environment, in which multiple seekers compete for the goal, is an interesting exercise.

course, implementing an algorithm that finds good values for these state-action pairs is the challenging part. Let's implement this idea of a Policy in first and worry about a suitable algorithm later.

*Example 3-7.*

**class Policy:**

```
def __init__(self, env):
    """A Policy suggest actions based on the current state.
    We do this by tracking the value of each state-action pair.
    """
    self.state_action_table = [
        [0 for _ in range(env.action_space.n)] for _ in range(env.observation_space.n) ❶
    ]
    self.action_space = env.action_space

def get_action(self, state, explore=True, epsilon=0.1):
    """Explore randomly or exploit the best value currently available."""
    if explore and random.uniform(0, 1) < epsilon: ❷
        return self.action_space.sample()
    return np.argmax(self.state_action_table[state]) ❸
```

- ❶ We define a nested list of values for each state-action pair, initialized to zero.
- ❷ Sometimes we might want to randomly explore actions in the game, which is why we introduce an explore parameter to the get\_action method. By default, this happens 10% of the time.
- ❸ We return the action with the highest value in the lookup table, given the current state.

I've sneaked in a little implementation detail into the Policy definition that might be a bit confusing. The get\_action method has an explore parameter. The reason for this is that if you learn an extremely poor policy, e.g. one that always wants you to move left, you have no chance of ever finding better solutions. In other words, sometimes you need to explore new ways, and not "exploit" your current understanding of the game. As indicated before, we haven't discussed how to learn to improve the values in the state\_action\_table of our policy. For now, just keep in mind that the policy gives us the actions we want to follow when simulating the maze game.

Moving on to the Simulation class we spoke about earlier, a simulation should take an Environment and compute actions of a given Policy until the goal is reached and the game ends. The data we observe when "rolling out" a full game like this is what we call the *experience* we gained. Accordingly, our Simulation class has a rollout



method that computes experiences for a full game and returns them. Here's what the implementation of the `Simulation` class looks like:

*Example 3-8.*

```
class Simulation(object):
    def __init__(self, env):
        """Simulates rollouts of an environment, given a policy to follow."""
        self.env = env

    def rollout(self, policy, render=False, explore=True, epsilon=0.1): ❶
        """Returns experiences for a policy rollout."""
        experiences = []
        state = self.env.reset() ❷
        done = False
        while not done:
            action = policy.get_action(state, explore, epsilon) ❸
            next_state, reward, done, info = self.env.step(action) ❹
            experiences.append([state, action, reward, next_state]) ❺
            state = next_state
            if render: ❻
                time.sleep(0.05)
                self.env.render()

        return experiences
```

- ❶ We compute a game “roll-out” by following the actions of a policy, and we can optionally render the simulation.
- ❷ To be sure, we reset the environment before each rollout.
- ❸ The passed in policy drives the actions we take. The `explore` and `epsilon` parameters are passed through.
- ❹ We step through the environment by applying the policy's action.
- ❺ We define an experience as a (`state`, `action`, `reward`, `next_state`) quadruple.
- ❻ Optionally render the environment at each step.

Note that each entry of the experiences we collect in a rollout consists of four values: the current state, the action taken, the reward received, and the next state. The algorithm we're going to implement in a moment will use these experiences to learn from them. Other algorithms might use other experience values, but those are the ones we need to proceed.

Now we have a policy that hasn't learned anything just yet, but we can already test its interface to see if it works. Let's try it out by initializing a `Simulation` object, calling its `rollout` method on a not-so-smart `Policy`, and then printing the `state_action_table` of it:

*Example 3-9.*

```
untrained_policy = Policy(environment)
sim = Simulation(environment)

exp = sim.rollout(untrained_policy, render=True, epsilon=1.0) ❶
for row in untrained_policy.state_action_table:
    print(row) ❷
```

❶ We roll-out one full game with an “untrained” policy that we render.

❷ The state-action values are currently all zero.

If you feel like we haven't made much progress since the last section, I can promise you that things will come together in the next one. The prep work of setting up a `Simulation` and a `Policy` were necessary to frame the problem correctly. Now the only thing that's left is to devise a smart way to update the internal state of the `Policy` based on the experiences we've collected, so that it actually learns to play the maze game.

## Training a Reinforcement Learning Model

Imagine we have a set of experiences that we've collected from a couple of games. What would be a smart way to update the values in the `state_action_table` of our `Policy`? Here's one idea. Let's say you're sitting at position (3,5), and you've decided to go right, which puts you at (4,5), just one step away from the goal. Clearly you could then just go right and collect a reward of 1 in this scenario. That must mean the current state you're in combined with an action of going “right” should have a high value. In other words, the value of this particular state-action pair should be high. In contrast, moving left in the same situation does not lead to anything, and the corresponding state-action pair should have a low value.

More generally, let's say you were in a given state, you've decided to take an action, leading to a reward, and you're then in `next_state`. Remember that this is how we defined an experience. With our `policy.state_action_table` we can peek a little ahead and see if we can expect to gain anything from actions taken from `next_state`. That is, we can compute

```
next_max = np.max(policy.state_action_table[next_state])
```

How should we compare the knowledge of this value to the current state-action value, which is `value = policy.state_action_table[state][action]`? There are many ways to go about this, but we clearly can't completely discard the current value and put too much trust in `next_max`. After all, this is just a single piece of experience we're using here. So as a first approximation, why don't we simply compute a weighted sum of the old and the expected value and go with `new_value = 0.9 * value + 0.1 * next_max`? Here, the values 0.9 and 0.1 have been chosen somewhat arbitrarily, the only important piece is that the first value is high enough to reflect our preference to keep the old value, and that both weights sum to 1. That formula is a good starting point, but the problem is that we're not at all factoring in the crucial information that we're getting from the reward. In fact, we should put more trust in the current reward value than in the projected `next_max` value, so it's a good idea to discount the latter a little, let's say by 10%. Updating the state-action value would then look like this:

```
new_value = 0.9 * value + 0.1 * (reward + 0.9 * next_max)
```

Depending on your level of experience with this kind of reasoning, the last few paragraphs might be quite a lot to digest. If you didn't quite get this last update rule for `new_value`, maybe consider coming back to it on a second read through. The good thing is that, if you've understood the explanations up to this point, the remainder of this chapter will likely come easy to you. Mathematically, this was the last (and only) hard part of this example<sup>2</sup>.

We're almost there, so let's formalize this procedure by implementing an `update_policy` function for a policy and collected experiences:

*Example 3-10.*

```
import numpy as np
```

```
def update_policy(policy, experiences):
    """Updates a given policy with a list of (state, action, reward, state) experiences."""
    weight = 0.1
    discount_factor = 0.9
    for state, action, reward, next_state in experiences: ❶
        next_max = np.max(policy.state_action_table[next_state]) ❷
        value = policy.state_action_table[state][action] ❸
        new_value = (1 - weight) * value + weight * (reward + discount_factor * next_max) ❹
        policy.state_action_table[state][action] = new_value ❺
```

2 If you've worked with RL before, you will have noticed by now that this is an implementation of the so-called Q-Learning algorithm. It's called like that, because the state-action table can be described as a function `Q(state, action)` that returns values for these pairs.

- ❶ We loop through all experiences in order.
- ❷ Then we choose the maximum value among all possible actions in the next state.
- ❸ We then extract the current state-action value.
- ❹ The new value is the weighted sum of the old value and the expected value, which is the sum of the current reward and the discounted next\_max.
- ❺ After updating, we set the new state\_action\_table value.

Having this function in place now makes it really simple to train a policy to make better decisions. We can use the following procedure:

- Initialize a policy and a simulation.
- Run the simulation many times, let's say for a total of 10000 runs.
- For each game, first collect the experiences by running a rollout.
- Then update the policy by calling update\_policy on the collected experiences.

That's it! The following train\_policy function implements the above procedure straight up.

*Example 3-11.*

```
def train_policy(env, num_episodes=10000):
    """Training a policy by updating it with rollout experiences."""
    policy = Policy(env)
    sim = Simulation(env)
    for _ in range(num_episodes):
        experiences = sim.rollout(policy) ❶
        update_policy(policy, experiences) ❷
    return policy
```

```
trained_policy = train_policy(environment) ❸
```

- ❶ Collect experiences for each game.
- ❷ Update our policy with those experiences.
- ❸ Finally, train and return a policy for our environment from before.

Note that the high-brow way of speaking of a full play-through of the maze game is an *episode* in the RL literature. That's why we call the argument `num_episodes` in the `train_policy` function, rather than `num_games`.

Now that we have a trained policy, let's see how well it performs. We've run random policies twice before in this chapter, just to get an idea of how well they work for the maze problem. But let's now properly evaluate our trained policy on several games and see how it does on average. Specifically, we'll run our simulation for a couple of episodes and count how many steps it took per episode to reach the goal. So, let's implement an `evaluate_policy` function that does precisely that:

*Example 3-12.*

```
def evaluate_policy(env, policy, num_episodes=10):
    """Evaluate a trained policy through rollouts."""
    simulation = Simulation(env)
    steps = 0

    for _ in range(num_episodes):
        experiences = simulation.rollout(policy, render=True, explore=False) ❶
        steps += len(experiences) ❷

    print(f"{steps / num_episodes} steps on average "
          f"for a total of {num_episodes} episodes.")
```

```
evaluate_policy(environment, trained_policy)
```

- ❶ This time we set `explore` to `False` to fully exploit the learnings of the trained policy.
- ❷ The length of the experiences is the number of steps we took to finish the game.

Apart from seeing the trained policy crush the maze problem ten times in a row, as we hoped it would, you should also see the following prompt:

```
8.0 steps on average for a total of 10 episodes.
```

In other words, the trained policy is able to find optimal solutions for the maze game. That means you've successfully implemented your first RL algorithm from scratch!

With the understanding you've built up by now, do you think placing the seeker into randomized starting positions and then running this evaluation function would still work? Why don't you go ahead and make the changes necessary for that?

Another interesting question to ask yourself is what assumptions went into the algorithm we used. For instance, it's clearly a prerequisite for the algorithm that all state-

action pairs can be tabulated. Do you think this would still work well if we had millions of states and thousands of actions?

## Building a Distributed Ray Trainer

Let's take a step back here. If you're an RL expert, you'll know what we've been doing the whole time. If you're completely new to RL, you might just be a little overwhelmed. If you're somewhere in between, you hopefully like the example but might be wondering how what we've done so far relates to Ray. That's a great question. As you'll see shortly, all we need to make the above RL experiment a distributed Ray app is writing three short code snippets. This is what we're going to do:

- We create a Ray task that can initialize a Policy remotely.
- Then we make the Simulation a Ray actor in just a few lines of code.
- After that we wrap the update\_policy function in a Ray task.
- Finally, we define a parallel version of train\_policy that's structurally identical to its original version.

Let's tackle the first two steps of this plan by implementing a create\_policy task and a Ray actor called SimulationActor:

*Example 3-13.*

```
import ray

ray.init()
environment = Environment()
env_ref = ray.put(environment) ❶

@ray.remote
def create_policy():
    env = ray.get(env_ref)
    return Policy(env) ❷

@ray.remote
class SimulationActor(Simulation): ❸
    """Ray actor for a Simulation."""
    def __init__(self):
        env = ray.get(env_ref)
        super().__init__(env)
```

- ❶ After initializing it, we put our environment into the Ray object store.
- ❷ This remote task returns a new Policy object.

- ③ This Ray actor wraps our `Simulation` class in a straightforward way.

With the foundations developed in chapter [Chapter 2](#) you should have no problems reading this code. It might take some getting used to writing it yourself, but conceptually you should be on top of this example.

Moving on, let's define a distributed `update_policy_task` Ray task and then wrap everything (two tasks and one actor) in a `train_policy_parallel` function that distributes this RL workload on your local Ray cluster:

*Example 3-14.*

```
@ray.remote
def update_policy_task(policy_ref, experiences_list):
    """Remote Ray task for updating a policy with experiences in parallel."""
    [update_policy(policy_ref, ray.get(xp)) for xp in experiences_list] ❶
    return policy_ref

def train_policy_parallel(num_episodes=1000, num_simulations=4):
    """Parallel policy training function."""
    policy = create_policy.remote() ❷
    simulations = [SimulationActor.remote() for _ in range(num_simulations)] ❸

    for _ in range(num_episodes):
        experiences = [sim.rollout.remote(policy) for sim in simulations] ❹
        policy = update_policy_task.remote(policy, experiences) ❺

    return ray.get(policy) ❻
```

- ❶ This task defers to the original `update_policy` function by passing a reference to a policy and experiences retrieved from the object store.
- ❷ To train in parallel, we first create a policy remotely, which returns a reference we call `policy`.
- ❸ Instead of one simulation, we create four simulation actors.
- ❹ Experiences now get collected from remote roll-outs on simulation actors.
- ❺ Then we can update our policy remotely. Note that `experiences` is a nested list of experiences.
- ❻ Finally, we return the trained policy by retrieving it from the object store again.

This allows us to take the last step and run the training procedure in parallel and then evaluate the resulting as before.

*Example 3-15.*

```
parallel_policy = train_policy_parallel()  
evaluate_policy(environment, parallel_policy)
```

The result of those two lines is the same as before, when we ran the serial version of the RL training for the maze. I hope you appreciate how `train_policy_parallel` has the exact same high-level structure as `train_policy`. It's a good exercise to compare the two line-by-line. Essentially, all it took to parallelize the training process was to use the `ray.remote` decorator three times in a suitable way. Of course, you need some experience to get this right. But notice how little time we spent on thinking about distributed computing, and how much time we could spend on the actual application code. We didn't need to adopt an entirely new programming paradigm and could simply approach the problem in the most natural way. Ultimately, that's what you want — and Ray is great at giving you this kind of flexibility.

To wrap things up, let's have a quick look at the task execution graph of the Ray application that we've just built. To recap, what we did was:

- The `train_policy_parallel` function creates several `SimulationActor` actors and a policy with `create_policy`
- The simulation actors create roll-outs with the policy and thereby collect experiences that `update_policy_task` uses to update the policy.
- This works, because of the way updating the policy is designed. It doesn't matter if the experiences were collected by one or multiple simulations.
- The rolling out and updating continues until we reached the number of episodes we want to train for, then the final `trained_policy` is returned.

Figure [Figure 3-1](#) summarizes this task graph in a compact way:



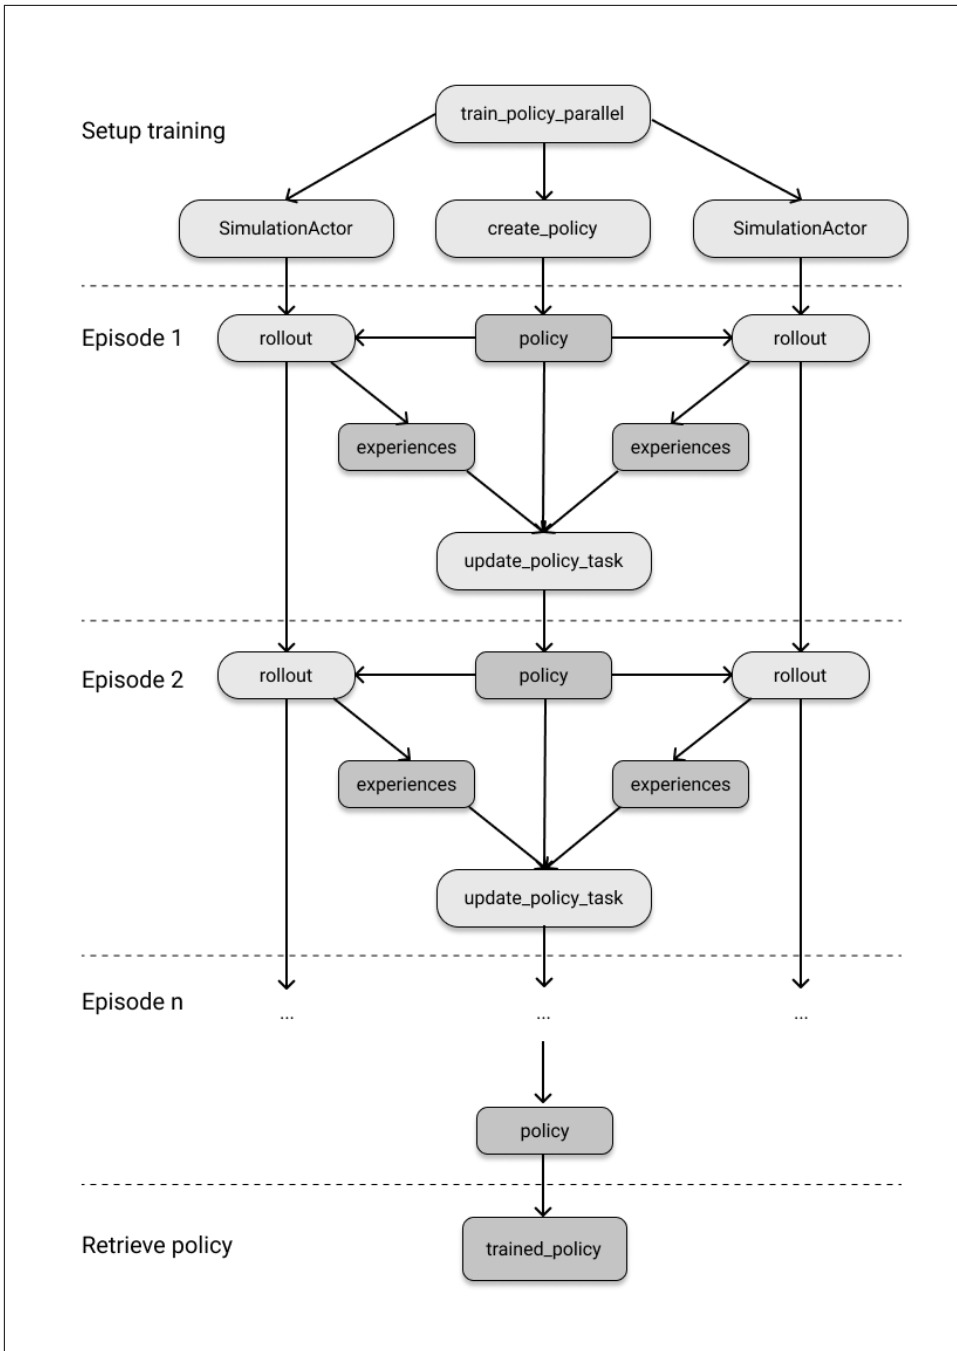


Figure 3-1. Parallel training of a reinforcement learning policy with Ray.

An interesting side note about the running example of this chapter is that it's an implementation of the pseudo-code example used to illustrate the flexibility of Ray in [the initial paper](#) by its creators. That paper has a figure similar to [Figure 3-1](#) and is worth reading for context.

## Summary

To recap, we've implemented a simple maze problem in plain Python and then solved the task of finding the goal in that maze using a straightforward reinforcement learning algorithm. We then took this solution and ported it to a distributed Ray application in roughly 25 lines of code. We did so without having to plan how to work with Ray — we simply used the Ray API to parallelize our Python code. This example shows how Ray gets out of your way and lets you focus on your application code. It also demonstrates how custom workloads that use advanced techniques like RL can be efficiently implemented and distributed with Ray.

In the next chapter, you'll build on what you've learned here and see how easy it is to solve our maze problem directly with the higher-level Ray RLlib library.

## About the Author

---

**Max Pumperla** is a data science professor and software engineer located in Hamburg, Germany. He's an active open source contributor, maintainer of several Python packages, author of machine learning books and speaker at international conferences. As head of product research at Pathmind Inc. he's developing reinforcement learning solutions for industrial applications at scale using Ray. Pathmind works closely with the AnyScale team and is a power user of Ray's RLlib, Tune and Serve libraries. Max has been a core developer of DL4J at Skymind, helped grow and extend the Keras ecosystem, and is a Hyperopt maintainer.

**Edward Oakes**

**Richard Liaw**